

The Chatman Equation:

A Receipt Calculus for Reality-Addressed Agentic AI

Sean Chatman
Independent Researcher

Genesis — Foundations, Paper 00

Abstract

This paper introduces the Chatman Equation, $\mathcal{A} = \mu(\mathcal{O}^*)$ with $\mathcal{R} = \text{receipt}(\mathcal{A})$, as a receipt calculus for agentic AI systems. The equation separates raw observation from admitted observation: an action is not manufactured from arbitrary model output, logs, claims, or tool responses, but only from an admitted subspace produced by finite, decidable obligations. This boundary is motivated by Rice’s theorem: arbitrary semantic truth cannot be decided by a total program, so execution must proceed through a restricted admission language rather than through unbounded interpretation.

The paper formalizes five objects: observation space, admitted space, manufacturing morphism, action/artifact space, and receipt space. It defines admission as a computable retraction, refusal as a first-class value, refusal composition as a denial monoid, and lifecycle governance as a typestate category whose illegal transitions are uninhabited. Manufacture is modeled as a deterministic bounded morphism from admitted observations to artifacts, while receipts are collision-committed projections of execution traces sized for bounded verification.

The enforced form, the Bounded Receipted Chatman Equation, requires four invariants: no actuation from raw or refused observations, bounded deterministic manufacture, total receipt emission, and conformance replay. Under these invariants, the paper proves conservation of consequence: every standing action has an admitted cause and a receipt-chain position, while receiptless actuation is excluded by construction. The result is a formal foundation for reality-addressed agentic systems in which scalable trust is obtained through admission, manufacture, receipt, and replay rather than through semantic omniscience or post hoc explanation.

Contents

Abstract	i
1 Introduction: One Equation, Five Objects	1
1.1 The claim in one line	1
1.2 Why comprehension is the wrong standard, in one paragraph	2
1.3 What this paper proves	2
2 The Computability Boundary: Rice Quarantine	3
2.1 The observation axiom	3
3 Admission Algebra: Refusal as Data	5
3.1 The admission map and refusal as data	5
3.2 The denial monoid and the eight-category taxonomy	6
4 The BRCE Byte: Machine Algebra of Standing	8
4.1 CPU notation for the BRCE byte	8
4.2 The machine algebra of lawful action	8
5 The Lifecycle as a Typestate Category	10
5.1 Stages, transitions, and the free path category	10
5.2 Illegal transition is a type error	11
5.3 Judgment carries its refusal	12
5.4 The 3-Node SEQ Lifecycle Token Game	12
6 Manufacture, Receipt, and the Enforced Equation	13
6.1 The manufacturing morphism	13
6.2 Receipts	13
6.3 The Bounded Receipted Chatman Equation	14
7 A Worked Example: A Contract-Claim Law Object	16
7.1 Raw	16
7.2 Judgment and an andon halt	16
7.3 Evidence, re-judgment, admission	17
7.4 Manufacture and receipt	17
8 Map of the Series	18

<i>CONTENTS</i>	iii
9 Conclusion	20
A Notation	21
B Code Correspondence	22

Chapter 1

Introduction: One Equation, Five Objects

1.1 The claim in one line

A prior paper in this program [1] advanced a single governing law, the *Chatman equation*

$$\mathcal{A} = \mu(\mathcal{O}), \tag{1.1}$$

read there as: an action state $\mathcal{A}$ is the deterministic image, under a measurement function μ , of a typed knowledge graph $\mathcal{O}$. That paper demonstrated the law at enterprise scale — knowledge hooks over RDF/SHACL, the KNHK hot-path executor, the **ggen** bounded projection, and Lockchain provenance — and reported operational constants (sub-two-nanosecond rule checks, full coverage of the forty-three workflow patterns of van der Aalst et al. [4], cryptographically receipted runs). It established that deterministic projection of typed knowledge into verifiable action is not a simulation but a deployed reality.

This paper does not restate that result; it *grounds* it. Equation (1.1) silently assumes its input is already typed — already lawful. The engineering that makes $\mathcal{O}$ typed, that decides which raw records may enter μ at all, and that emits the receipts by which a bounded human trusts the output, is exactly the machinery this series formalizes. We therefore refine (1.1) by *factoring* the projection into an admission stage and a manufacturing stage, exposing the admitted sub-space $\mathcal{O}^*$ and the receipt space $\mathcal{R}$ that (1.1) left implicit. The refined identity, which governs the whole series, is

$$\boxed{\mathcal{A} = \mu(\mathcal{O}^*)} \quad \text{with} \quad \mathcal{O}^* = \text{im}(\alpha), \quad \alpha : \mathcal{O} \rightarrow \mathcal{O}^* \cup \{\perp\}. \tag{1.2}$$

In words: nothing is manufactured from a raw observation; only from an *admitted* one; and every manufacture leaves a receipt. The five objects of the equation are named once, here, and used unchanged throughout the series.

Definition 1.1 (The five objects). The Chatman equation (1.2) relates five objects.

- (i) $\mathcal{O}$, the *observation space*: arbitrary finite records — logs, model outputs, sensor frames, human claims, tool responses — carrying no decidable semantics (Axiom 2.1).
- (ii) $\mathcal{O}^* = \mathcal{O}^* \subseteq \mathcal{O}$, the *admitted space*: the decidable sub-collection onto which a finite battery of obligations is computable, together with the distinguished refusal $\perp$ (Definition 3.1).

- (iii) $\mu : \mathcal{O}^* \rightarrow \mathcal{A}$, the *manufacturing morphism*: a deterministic, bounded, computable map from admitted observations to artifacts (Definition 6.1).
- (iv) $\mathcal{A}$, the *artifact/action space*: the actuated outputs, exhausted by $\text{im } \mu$ (Definition 6.1).
- (v) $\mathcal{R}$, the *receipt space*: bounded, collision-committed projections of the execution that produced an artifact, sized for a human verifier (Definition 6.2).

Remark (Relation to the prior μ). The measurement function of (1.1) is the composite of this paper's two stages: $\mu_{\text{prior}} = \mu \circ \alpha$, with α the admission of Definition 3.1 and μ the manufacture of Definition 6.1. Nothing in [1] is retracted; the ingress guards H of that paper *are* the obligation battery, and its cryptographic receipts *are* the objects of $\mathcal{R}$. This series simply refuses to leave α and $\mathcal{R}$ implicit.

1.2 Why comprehension is the wrong standard, in one paragraph

The systems this program serves exceed any human's working-memory capacity, a fixed small integer we write κ (canonically $\kappa \approx 4$ following Miller and Cowan [7, 8]). A system whose faithful description needs more than κ chunks cannot be trusted by comprehension. The response of this series is not to shrink the system but to insert, between it and the human, a map of codomain dimension $\leq \kappa$ whose fidelity is guaranteed by mechanism rather than by the reader — the receipt. That is the subject of the keystone paper *The Projection Principle*; here we build the algebraic and type-theoretic floor beneath it.

1.3 What this paper proves

Chapter 2 axiomatizes observation and proves the Rice specialization that forces admission to be a retraction, not a decision; it defines the admission map, refusal, and the denial monoid, and states the totality property of the eight-category refusal taxonomy. Chapter 5 builds the lifecycle typestate category and proves the illegal-transition-is-a-type-error theorem. Chapter 6 defines the manufacturing morphism and the receipt, states the Chatman equation as a factoring, and gives the enforced form BRCE with its conservation corollary. Chapter 7 works a contract-claim law object end to end. Chapter 8 is the map of the series. Throughout, a claim is a definition, a theorem with proof, or a citation; there is no fourth kind of sentence.

Chapter 2

The Computability Boundary: Rice Quarantine

2.1 The observation axiom

We begin with the least structure that makes “lawful action” expressible.

Axiom 2.1 (Observation). There is a set $\mathcal{O}$, the observation space, whose elements are arbitrary finite records. $\mathcal{O}$ carries no decidable semantics: the predicate “does this observation mean what it purports to mean” is not assumed computable. Concretely, an element of $\mathcal{O}$ is any finite byte string a caller may submit; in the **praxis** implementation it is the JSON payload handed to the `law judge` verb, before any obligation has been evaluated.

Axiom 2.1 is not a limitation we impose; it is forced by a classical result, which we specialize to observations.

Theorem 2.1 (Rice, specialized to observations). *Let $\mathcal{U}$ be a universal model of computation and let observations in $\mathcal{O}$ range over finite encodings that may denote arbitrary $\mathcal{U}$ -programs (equivalently, the partial functions they compute). Let P be any non-trivial semantic property of those denoted functions — non-trivial meaning some observation has it and some does not, and P depends only on the denoted function, not on the encoding. Then the set $\{o \in \mathcal{O} : P(o)\}$ is undecidable.*

Proof. This is Rice’s theorem [2] instantiated at $\mathcal{O}$. Rice’s theorem states that every non-trivial property of the partial function computed by a Turing machine is undecidable. Since observations are unrestricted finite encodings, the map $e : \mathcal{O} \rightarrow \{\mathcal{U}\text{-programs}\}$ sending an observation to the program it denotes is onto a set closed under the standard constructions, and P factors through the denoted function by hypothesis. Suppose a decider D existed for $\{o : P(o)\}$. Fix observations o_+ with P and o_- without (non-triviality). Given any machine M and input x , build the observation $o_{M,x}$ denoting the program “run $M(x)$; if it halts, behave as $e(o_+)$; else diverge.” Then $o_{M,x}$ denotes a function with property P iff $M(x)$ halts, so $D(o_{M,x})$ decides the halting problem, a contradiction. Hence no such D exists. $\square$

Corollary 2.2 (Admission cannot be semantic decision). *There is no algorithm that admits an observation $o \in \mathcal{O}$ by deciding a non-trivial semantic property of what o denotes. Any total,*

terminating admission procedure must therefore decide a syntactic, decidable surrogate — it retracts $\mathcal{O}$ onto a decidable sub-language rather than deciding meaning.

Proof. Immediate from Theorem 2.1: a total admission that decided a non-trivial semantic property would be a decider for an undecidable set. A total terminating procedure exists only for decidable predicates; the largest such structure available is a decidable sub-collection, onto which admission maps. $\square$

Corollary 2.2 is the reason the entire apparatus exists. One does not admit an observation by understanding it. One admits it by *retracting* it onto a sub-language — the *Rice quarantine* — on which the properties of interest are, by construction, decidable.

Chapter 3

Admission Algebra: Refusal as Data

3.1 The admission map and refusal as data

Definition 3.1 (Admitted space, obligations, admission). Fix a decidable sub-collection $\mathcal{O}^* \subseteq \mathcal{O}$ and a finite battery of *obligations* $g_1, \dots, g_m : \mathcal{O} \rightarrow \{0, 1\}$, each total and computable. The *admission map* is the partial retraction

$$\alpha : \mathcal{O} \rightarrow \mathcal{O}^* \cup \{\perp\}, \quad \alpha(o) = \begin{cases} \rho(o) \in \mathcal{O}^*, & \text{if } \bigwedge_{i=1}^m g_i(o) = 1, \\ \perp, & \text{otherwise,} \end{cases}$$

where $\rho : \mathcal{O} \rightarrow \mathcal{O}^*$ is a computable normalization fixing $\mathcal{O}^*$ pointwise ($\rho|_{\mathcal{O}^*} = \text{id}_{\mathcal{O}^*}$) and $\perp$ is the distinguished *refusal*. Admission is idempotent on its image: $\alpha \circ \alpha = \alpha$.

Proposition 3.1 (Admission is a retraction). *α restricted to $\mathcal{O}^*$ is the identity, and $\alpha \circ \alpha = \alpha$ as partial maps. Hence $\mathcal{O}^*$ is a retract of $\text{dom}(\alpha)$, and re-admitting an already-admitted observation neither changes it nor can newly refuse it.*

Proof. For $x \in \mathcal{O}^*$ every $g_i(x) = 1$ (obligations are decidable on the admitted language by construction), so $\alpha(x) = \rho(x) = x$. Thus $\alpha|_{\mathcal{O}^*} = \text{id}_{\mathcal{O}^*}$. For general $o \in \text{dom}(\alpha)$ with $\alpha(o) \neq \perp$ we have $\alpha(o) \in \mathcal{O}^*$, whence $\alpha(\alpha(o)) = \alpha(o)$; and $\alpha(\perp) = \perp$ by convention. Idempotence follows on all of $\text{dom}(\alpha)$. $\square$

Remark (Refusal is a value, not an exception). $\perp$ is a first-class output of a correctly functioning admission: the system computed that an obligation failed and recorded *which*. In **praxis** this is the type **Andon**, whose **Halted** variant carries both the unmet obligations and their refusal classification:

```
Andon::Halted{ unmet: Vec<Obligation>, refusals: Vec<RefusalScenario>, at: u64 }.
```

A refusal is data with a timestamp and a cause, not a thrown error. This is the formal content of the program's slogan *refusal is data*.

The obligations are themselves first-class, hashable values, not opaque closures. The **praxis Obligation** enum realizes exactly three decidable kinds, matching the three ways an observation can fail to be admitted:

Obligation variant	meaning	g_i decides
<code>Precondition{predicate_id, params_hash}</code>	a named predicate must hold	predicate membership
<code>BlockingConstraint{reason}</code>	a hard block until lifted	a flag
<code>EvidenceRequired{evidence_type}</code>	external evidence must be present	presence of a typed record

Each is a syntactic, terminating check — a legitimate g_i under Corollary 2.2 — never a semantic decision about what the payload “really means.”

3.2 The denial monoid and the eight-category taxonomy

Obligation checking composes, and the composition carries the algebra used throughout the series’ geometry.

Construction 3.1 (Denial monoid). Let $D = (\{0, 1\}^n, \vee, \mathbf{0})$ be the commutative idempotent monoid of *denial words*: n independent lanes, componentwise OR, identity $\mathbf{0}$ (“admitted, all lanes clear”). Each refusal cause contributes a lane map $d_i : \mathcal{O} \rightarrow \{0, 1\}^n$ with $d_i(o) = \mathbf{0} \iff$ cause i is absent. The *total denial* of an observation is $d(o) = \bigvee_i d_i(o)$, and $\alpha(o) \neq \perp \iff d(o) = \mathbf{0}$.

Proposition 3.2 (D is a bounded join-semilattice; denial is monotone). $(\{0, 1\}^n, \vee, \mathbf{0})$ is a commutative idempotent monoid — the join-semilattice of the Boolean lattice $2^{[n]}$ with bottom $\mathbf{0}$. Consequently adding obligations can only enlarge the denial (turn admits into refusals), never the reverse.

Proof. Idempotence $x \vee x = x$, commutativity, and associativity hold coordinatewise; a product of semilattices is a semilattice, with $\mathbf{0}$ the identity. For monotonicity, if $G' \supseteq G$ are cause sets then $d_{G'}(o) = d_G(o) \vee \bigvee_{i \in G' \setminus G} d_i(o) \succeq d_G(o)$ coordinatewise, and $d(o) = \mathbf{0}$ is the unique minimum, so $\{o : d_{G'}(o) = \mathbf{0}\} \subseteq \{o : d_G(o) = \mathbf{0}\}$. $\square$

This construction is realized in code. The `bcinr_powl_receipt::denial::DenialPolarity` type is a newtype over `u64` with a zero element `ADMITTED` and seven named non-zero lanes; its `compose` operation is bitwise OR; and `praxis`’s `compose.denials` folds a set of refusal scenarios through `compose` starting from `ADMITTED` — literally the monoid fold $\bigvee_i d_i$ with the identity as the empty product. A refusal *category* is then the coarse classification of a lane.

Definition 3.2 (Refusal taxonomy). A *refusal taxonomy* is a pair (S, cat) where S is a finite set of concrete refusal scenarios and $\text{cat} : S \rightarrow C$ is a total map onto a fixed finite set C of categories. In `praxis`, C is the eight-element set

$$C = \{\text{Identity, Capacity, Topology, Temporal, Lifecycle, Authorization, Prerequisites, Reserved}\},$$

S is the `RefusalScenario` enum (obligation-driven halts, the seven `DenialPolarity` lanes, and the `prolog8` kernel / andon-ring denials), and `cat = RefusalScenario::category`.

Proposition 3.3 (Totality is mechanized). The category map $\text{cat} : S \rightarrow C$ is total, and its totality is enforced by the host type system: `RefusalScenario::category` is a *match* with no wildcard arm, so a new scenario variant added without a category is a compile error. Moreover the lane assignment $\text{lane} : S \rightarrow D$ (`denial_lane`) and the single-lane inverse $\text{scn} : D \rightarrow S$ (`scenario_for_denial_lane`) satisfy $\text{lane} \circ \text{scn} = \text{id}$ on each of the seven named lanes, i.e. scn is a section of lane over the single-lane words.

Proof. Totality of a wildcard-free exhaustive `match` over a closed enum is exactly the compiler’s exhaustiveness guarantee: the program does not type-check unless every variant is covered, so `category` denotes a total function $S \rightarrow C$. The section identity is the round-trip property verified by the module’s test `category_mapping_is_total_over_denial_lanes`: for each of the seven single-lane constants ℓ , `scn`(ℓ) is defined and `lane`(`scn`(ℓ)) = ℓ . That `scn` is only a *partial* section (undefined on composed multi-lane words and on `ADMITTED`) reflects that a composed denial has no single scenario; it is a join $\bigvee$ of several, per Construction 3.1. $\square$

Proposition 3.3 is the algebraic backbone of the admission pillar (Chapter 8): the taxonomy is not documentation but a total function whose totality the compiler certifies, and the denial lattice is the semantic domain of that function.

Chapter 4

The BRCE Byte: Machine Algebra of Standing

4.1 CPU notation for the BRCE byte

When O^* is projected into the machine's fastest lawful state, it becomes a single byte. We now assign notation to the physical operations by which the Bounded Receipted Chatman Equation (BRCE) achieves execution at the picosecond scale.

Let $\mathbb{B} = \{0, 1\}$ and let the BRCE standing byte be $b \in \mathbb{B}^8$. Write $0 = 00000000_2$ and $1 = 11111111_2$. Each basis bit is $e_i \in \mathbb{B}^8$, for $0 \leq i < 8$, where e_i has a 1 in lane i and 0 everywhere else. A standing byte is $b = b_0e_0 \vee b_1e_1 \vee \cdots \vee b_7e_7$. In the canonical instance, the lanes represent:

$$\begin{array}{llll} b_0 = \text{admitted}, & b_1 = \text{evidenced}, & b_2 = \text{budgeted}, & b_3 = \text{authorized}, \\ b_4 = \text{healthy}, & b_5 = \text{conformant}, & b_6 = \text{receipted}, & b_7 = \text{replayable}. \end{array}$$

Define the register-level CPU operation alphabet $\mathcal{I}_8$:

$$\mathcal{I}_8 = \{\wedge_8, \vee_8, \oplus_8, \neg_8, \ll_k, \gg_k, =_c, z, nz, \text{sel}, \text{pop}_8\}.$$

These correspond to the primitive machine instructions acting directly on fixed-width state: bitwise AND ($\wedge_8$), OR ($\vee_8$), XOR ($\oplus_8$), NOT ($\neg_8$), constant comparison ($=_c$), zero test (z), nonzero test (nz), branchless select (sel), and population count (pop_8). A lawful BRCE hot path is a word over $\mathcal{I}_8$.

4.2 The machine algebra of lawful action

With the CPU alphabet defined, every governance operation reduces to a bounded Boolean map over $\mathbb{B}^8$.

Admission and Denial. The denial byte is the complement of the standing byte: $D(b) = \neg_8 b$. No denial exists when $D(b) = 0$. Admission is exactly the zero-test of the denial byte:

$$A_{\text{gate}}(b) = z(D(b)) \implies A_{\text{gate}}(b) = 1 \iff b = 1.$$

Refusal is the nonzero test: $H(b) = \text{nz}(D(b))$. If $d^{(1)}, \dots, d^{(n)}$ are denial bytes from independent obligations, their composed denial is simply the bitwise OR: $d_\Sigma = d^{(1)} \vee_8 d^{(2)} \vee_8 \dots \vee_8 d^{(n)}$. Admission exists exactly when $d_\Sigma = 0$.

Policy masks and repairs. Let $m \in \mathbb{B}^8$ be a required policy mask. The policy predicate is $P_m(b) = [(b \wedge_8 m) = m]$, meaning every lane required by m is present in b . The missing lanes are isolated by $M_m(b) = m \wedge_8 \neg_8 b$, so $P_m(b) = 1 \iff M_m(b) = 0$. A policy is just a mask; a violation is simply “mask AND NOT byte”. Repairing a lane i is $b' = b \vee_8 e_i$.

Completeness and Branchless Selection. Completeness is the population count $C(b) = \text{pop}_8(b) \in \{0, \dots, 8\}$. Repair distance is $R(b) = 8 - \text{pop}_8(b)$. This allows the system to determine that an object is “missing one lane” without inspecting raw meaning. For a predicate bit $p \in \mathbb{B}$, define its full-byte mask $\hat{p} = 1$ if $p = 1$ else 0. Branchless select is $\text{sel}(p, u, v) = (\hat{p} \wedge_8 u) \vee_8 (\neg_8 \hat{p} \wedge_8 v)$.

Action selection. Define the byte-to-integer decoder $\nu(b) = \sum_{i=0}^7 b_i 2^i$, mapping $\mathbb{B}^8 \rightarrow \{0, \dots, 255\}$. Reserving zero ($\nu(b) = 0 \implies \perp$), the action selector is:

$$\sigma(b) = \begin{cases} \perp, & \nu(b) = 0, \\ a_{\nu(b)}, & 1 \leq \nu(b) \leq 255. \end{cases}$$

Thus $|\text{im}(\sigma) \setminus \{\perp\}| \leq 255$. Eight bits can select 255 non-null lawful actions without branching.

The constant-depth eligibility theorem. Let τ be the cycle quantum of the target machine. For primitive register operations, $\text{cost}(i) = 1$ for $i \in \mathcal{I}_8$. The execution depth $\text{depth}(f)$ of a composed expression f is its longest dependency chain in $\mathcal{I}_8$. The bounding constraint is $T(f) \approx \text{depth}(f) \cdot \tau$. The admission check $G(b) = [b = 1]$ has depth 1. The policy mask $P_m(b)$ has depth 2. From first principles, BRCE proves that admission, denial composition, policy masking, and action selection are finite Boolean maps with depth $\in \{1, 2, 3\}$. The hot-path eligibility checks compile to constant-depth Boolean expressions over fixed-width registers. On machines whose primitive register operations retire within a sub-nanosecond cycle budget, the admission predicate is picosecond-eligible.

Thus BRCE does not merely claim that eight bits are compact; it gives a CPU-level algebra in which lawful standing, denial, policy, repair, and action selection are all words over the machine’s primitive register operations. BRCE opens high-frequency domains by making law executable as byte algebra.

Chapter 5

The Lifecycle as a Typestate Category

5.1 Stages, transitions, and the free path category

An admitted observation does not become a receipt in one step. It traverses a fixed lifecycle, and the *illegality of skipping steps* is the first structural guarantee a newcomer must trust. We make it a theorem.

Definition 5.1 (Lifecycle category). Let **Life** be the category with four objects, the lifecycle stages

Raw, Validated, Admitted, Receipted,

and with non-identity morphisms generated by exactly three arrows

$$j : \text{Raw} \rightarrow \text{Validated}, \quad a : \text{Validated} \rightarrow \text{Admitted}, \quad r : \text{Admitted} \rightarrow \text{Receipted},$$

(*judge, admit, receipt*), closed under composition and identities. **Life** is the free category on the linear quiver $\text{Raw} \xrightarrow{j} \text{Validated} \xrightarrow{a} \text{Admitted} \xrightarrow{r} \text{Receipted}$.

Proposition 5.1 (Hom-sets of **Life**). *For stages X, Y , the hom-set $\mathbf{Life}(X, Y)$ has exactly one element if Y is reachable from X along the quiver (including $X = Y$, the identity) and is empty otherwise. In particular $\mathbf{Life}(\text{Raw}, \text{Receipted}) = \{r \circ a \circ j\}$ is a singleton, and there is no morphism $\text{Raw} \rightarrow \text{Admitted}$ that is not $a \circ j$, no morphism $\text{Validated} \rightarrow \text{Receipted}$ other than $r \circ a$, and no backward or skipping morphism at all.*

Proof. In the free category on a quiver, morphisms $X \rightarrow Y$ are directed paths from X to Y modulo nothing (paths are the morphisms). The quiver here is a simple directed path with no branches and no cycles, so between any two vertices there is at most one directed path: the sub-path of the unique maximal path, if Y lies after X , and none otherwise. Hence each nonempty hom-set is a singleton and each “illegal” hom-set (backward, or skipping an intermediate vertex) is empty. $\square$

5.2 Illegal transition is a type error

The `praxis` implementation represents a law object as a type carrying its stage as a phantom parameter:

```
LawObject<Payload, S: Stage, Law>,
```

where `Stage` is a *sealed* trait implemented only by the four zero-sized types `Raw`, `Validated`, `Admitted`, `Receipted`, and the phantom fields `_stage`, `_law` are private to the defining module. The three lifecycle arrows are the only stage-changing operations exposed: `Judge::judge` has type `Raw → Validated` (returning the halted `Raw` on refusal), `Admit::admit` has type `Validated → Admitted`, and `receipt` is a method available *only* on `LawObject<_, Admitted, _>`, with signature `Admitted → Receipted`, consuming its receiver. We now state the correspondence precisely.

Theorem 5.2 (Illegal lifecycle transitions are uninhabited). *Let $\mathcal{T}$ be the host type system and interpret each stage X as the type family `LawObject<P, X, L>` and each exposed stage-changing operation as a term of the corresponding function type. Then the set of stage-changing operations denotes exactly the generating arrows $\{j, a, r\}$ of **Life**, and consequently:*

- (i) *every well-typed stage-changing program is a composite of j, a, r , i.e. a morphism of **Life** (Proposition 5.1);*
- (ii) *for any pair (X, Y) with $\mathbf{Life}(X, Y) = \emptyset$ — a backward or step-skipping transition — there is no term of type `LawObject<P, X, L> → LawObject<P, Y, L>` constructible from the exposed interface; such a program does not type-check;*
- (iii) *no law object can be receipted twice: r consumes its `Admitted` receiver, so the term is linear in its argument and `Receipted` has no outgoing stage-changing arrow.*

Proof. (i) The only public constructor produces a `Raw` object; the only public stage-changing operations are `judge`, `admit`, `receipt`, whose types are j, a, r respectively. The stage transition helper that rebuilds an object at a new stage is `crate-private (pub(crate))`, and the phantom fields are private, so no external term can fabricate a `LawObject` at a chosen stage. Hence every externally well-typed stage-changing program is a composition of j, a, r , which by definition is a morphism of the free category **Life**. (ii) Suppose $\mathbf{Life}(X, Y) = \emptyset$. By Proposition 5.1 there is no path $X \rightarrow Y$ in the quiver, so no composite of j, a, r has type `LawObject<P, X, L> → LawObject<P, Y, L>`. Since (i) shows composites of j, a, r are the only such terms, the required term does not exist; a program asserting it fails to type-check. Concretely, `receipt` is impl'd only on the `Admitted` type, so calling it on `Raw` or `Validated` is a missing-method type error, and `admit` takes `Validated` by value, so feeding it a `Raw` is a type mismatch. (iii) `receipt` takes `self` by value (move), consuming the `Admitted` object; the returned `Receipted` type has no method producing a further stage. Thus the argument is used exactly once and the terminal object `Receipted` emits no stage-changing arrow, so a second receipting is both unconstructible (no arrow out of `Receipted`) and unavailable (the receiver was moved). $\square$

Corollary 5.3 (The admissible sub-category is what compiles). *The sub-category of “admissible morphisms” — those a program may actually perform — is precisely **Life**, and it coincides with the set of well-typed stage-changing programs over the `LawObject` interface. “Illegal transition is a type error” is therefore not a runtime check to be tested but a statement about which hom-sets are inhabited, discharged by the type checker before the program runs.*

Proof. Combine Theorem 5.2(i)–(ii): the inhabited stage-changing types are exactly the morphisms of **Life**, and the uninhabited ones are exactly the empty hom-sets. The admissible sub-category is thus **Life** itself, verified statically. $\square$

Remark (Where the parallel to admission lives). Chapter 2 retracted an undecidable observation space onto a decidable admitted language; here the type checker retracts the space of *all conceivable stage sequences* onto the decidable (indeed, trivially decidable) language of paths in **Life**. Both are Rice quarantines: an undecidable or unwieldy space of behaviors is replaced by a small decidable sub-language whose membership is mechanically certified. Admission does it at values; typestate does it at programs.

5.3 Judgment carries its refusal

The arrow $j : \text{Raw} \rightarrow \text{Validated}$ is total in a specific sense: it never throws, it either advances the object or returns it halted, carrying its denial. In *praxis*,

$\text{Judge}::\text{judge} : \text{LawObject}\langle P, \text{Raw}, L \rangle \longrightarrow \text{Result}\langle \text{LawObject}\langle P, \text{Validated}, L \rangle, \text{LawObject}\langle P, \text{Raw}, L \rangle \rangle,$

where the **Err** branch is the *same object*, now in **Andon**:**Halted** with its **unmet** obligations and **refusals**. This is Definition 3.1 at the type level: $\alpha(o)$ is either an element of $\text{Validated} \subseteq \mathcal{O}^*$ or the refusal $\perp$, and the refusal is a value one can inspect, not control flow one must catch. The admit arrow $a : \text{Validated} \rightarrow \text{Admitted}$ is analogous, returning **Err**(**Andon**) when admission is denied.

5.4 The 3-Node SEQ Lifecycle Token Game

The lifecycle transitions of the **LawObject** parameters represent a sequential path: **judge** $\rightarrow$ **admit** $\rightarrow$ **receipt**.

Definition 5.2 (3-Node SEQ POWL Token Game). The lifecycle process is a safe Petri net with state $M \in \{0, 1\}^4$ and transitions $T = \{j, a, r\}$ representing the judge, admit, and receipt actions. The token state is:

$$M = (\text{TOK_START}, \text{TOK_JUDGED}, \text{TOK_ADMITTED}, \text{TOK_DONE}), \quad (5.1)$$

with firing rules enabled by coordinatewise subtraction and addition:

$$M_{\text{next}} = (M \setminus \mathbf{m}^-) \cup \mathbf{m}^+.$$

Fitness $\varphi \in [0, 1]$ is computed in Q16.16 format. An out-of-order transition triggers a token shortfall and results in $\varphi < 1.0$.

Chapter 6

Manufacture, Receipt, and the Enforced Equation

6.1 The manufacturing morphism

Definition 6.1 (Manufacturing morphism). A *manufacturing morphism* is a deterministic computable map $\mu : \mathcal{O}^* \rightarrow \mathcal{A}$ subject to two laws:

- (M1) **Determinism / reproducibility:** $\mu(x)$ depends only on x ; equal admitted inputs yield bit-identical artifacts.
- (M2) **Boundedness:** μ factors through a representation of size bounded by fixed structural constants, so μ terminates with a priori bounded cost.

μ is undefined on $\mathcal{O} \setminus \mathcal{O}^*$ by construction, and refusal propagates: $\mu(\perp) = \perp$.

The governing identity of Definition 1.1, $\mathcal{A} = \mu(\mathcal{O}^*)$, is now precise: operationally $a = \mu(\alpha(o))$ whenever $\alpha(o) \neq \perp$, and undefined (a receipted refusal) otherwise. In **praxis**, one instance of μ is the manufacture of a PDDL8 planning domain and problem from a **pd1**: RDF ontology (the **mfg pd1** verb over **bcinr-pddl**): the ontology graph is hashed with BLAKE3 (a **graph_hash** witnessing determinism), and identical ontologies manufacture identical domain text, exactly (M1). The grounding-and-solving stage is (M2)’s bounded representation: a finite lifted domain grounded to a finite action set the planner searches.

Proposition 6.1 (Determinism makes μ a function on receipts). *Under (M1), the assignment $x \mapsto \mu(x)$ is a function (single-valued), hence the composite $o \mapsto \mu(\alpha(o))$ is a partial function on $\mathcal{O}$. Two runs on the same admitted input therefore agree bit-for-bit, and any disagreement is evidence that the input, not the morphism, differed.*

Proof. (M1) is exactly single-valuedness. Composition of the partial function α with the function μ is a partial function. Bit-identity of outputs on equal inputs is the contrapositive of the last clause. $\square$

6.2 Receipts

A manufacture leaves an *execution trace* $\sigma \in \Sigma$: every intermediate marking and sub-artifact. Traces live in a space of unbounded dimension relative to κ . The receipt is the bounded,

collision-committed projection of that trace.

Definition 6.2 (Receipt and chain). Fix a collision-resistant hash $H : \{0, 1\}^* \rightarrow \{0, 1\}^{256}$ (in *praxis*, BLAKE3 [5]). For a manufacture with payload bytes b , previous chain value h_- , and metadata θ (instruction index, activity, node kind, timestamp, denial word), the *frame* is a fixed-width encoding $\text{fr} = \langle \theta, H(b) \rangle$ and the *chained receipt* is

$$h_+ = H(h_- \parallel \text{fr}). \quad (6.1)$$

The *receipt* is the tuple $r = (\text{verdict}, h_+, \varphi, \text{reason}) \in \mathcal{R}$ of at most κ chunks: a Boolean verdict, the 256-bit chain value, a scalar conformance fitness φ , and a refusal reason (empty on accept).

Equation (6.1) is the `OcelCausalFrame` construction in *praxis*: the 32-byte BLAKE3 payload hash is packed into the frame’s eight object-reference words as little-endian `u32s` (repurposing the OCEL [6] object slots as a content commitment), the honest denial word from Chapter 2 is written into the frame’s `denial/fired_mask`, and the chain step is `chain_hash(t+1) = H(chain_hash(t) || frame_bytes(t+1))`. Determinism is tested directly: identical inputs produce identical chain hashes; differing payloads, previous hashes, instruction ids, or denial words produce differing chain hashes.

Definition 6.3 (Conformance fitness). Let $\varphi \in [0, 1]$ be one minus the fraction of tokens a replay was forced to consume on unenabled transitions of the lifecycle process model. $\varphi = 1$ iff the replay never attempted a disabled firing. In *praxis*, the lifecycle is the fixed three-node sequence `judge` $\rightarrow$ `admit` $\rightarrow$ `receipt`, replayed by `PowlReplayVerifier`; a lawful in-order trace yields fitness `0x0001_0000` — the value 1.0 in Q16.16 fixed point — while any out-of-order or step-skipping trace is a `TokenNotEnabled` violation with $\varphi < 1$.

The full theory of the receipt — that a bounded verifier reading $O(1)$ symbols rejects any perturbed trace except with negligible probability — is the subject of the receipt cryptography pillar (Chapter 8) and the keystone’s Faithful Projection Theorem. We state here only the commitment property we need for conservation.

Lemma 6.2 (Chain commitment). *Under the collision resistance of H , the chain value h_+ of (6.1) is a binding commitment to the pair (h_-, fr) : producing a distinct (h'_-, fr') with the same h_+ requires exhibiting a collision of H . By induction along the chain, h_+ binds the entire causal prefix.*

Proof. If $H(h_- \parallel \text{fr}) = H(h'_- \parallel \text{fr}')$ with distinct arguments, the two arguments are a collision of H , which occurs with negligible probability under collision resistance. Since h_- is itself such a commitment to the previous frame, an induction on chain length shows h_+ commits to every frame in the prefix. $\square$

6.3 The Bounded Receipted Chatman Equation

The Chatman equation $\mathcal{A} = \mu(\mathcal{O}^*)$ is a mathematical identity; its *enforced* form is the conjunction of invariants a running system maintains. We name it once for the series.

Definition 6.4 (BRCE: Bounded Receipted Chatman Equation). A system satisfies the *Bounded Receipted Chatman Equation* if for every actuated artifact $a \in \mathcal{A}$ it maintains all four invariants:

- (B1) **Admission gate:** $a = \mu(x)$ for some $x = \alpha(o) \neq \perp$; nothing is actuated from a raw or refused observation ($\text{im } \mu$ exhausts the actuated set).
- (B2) **Bounded manufacture:** μ satisfies (M1)–(M2); manufacture is deterministic and a priori cost-bounded.
- (B3) **Receipt totality:** every actuation appends exactly one chained receipt (6.1); there is no unreceipted actuation.
- (B4) **Conformance:** the actuation’s lifecycle replay has fitness $\varphi = 1$ (Definition 6.3); the trace stays within the process model.

Theorem 6.3 (Conservation of Consequence). *Under BRCE, the map $a \mapsto h_+(a)$ from actuated artifacts to receipt-chain positions is well defined and injective up to hash collision, and every actuated artifact is the image under μ of exactly one admitted observation. Equivalently: consequence is neither created without an admitted cause (B1) nor destroyed without a receipt (B3), and the chain preserves order.*

Proof. By (B1) actuation implies $a = \mu(x)$ with $x = \alpha(o) \neq \perp$, so $\text{im } \mu$ is the whole actuated set and its complement is never actuated. By (B2)/(M1) and Proposition 6.1, $x \mapsto \mu(x)$ is single-valued, so each actuated a has a determined admitted preimage x (unique as the input that produced it under a deterministic μ ; distinct admitted inputs yielding the same artifact are identified by the artifact, which is all conservation asserts). By (B3) the actuation appends exactly one chain position, so $a \mapsto h_+(a)$ is a well-defined total map on actuated artifacts. Injectivity up to collision is Lemma 6.2: if two distinct actuations shared a chain value, their frames would collide under H . Order preservation is the recursive dependence of h_+ on h_- in (6.1). $\square$

Corollary 6.4 (The antibody clause). *Define the overclaim set as actuations lacking a valid receipt. Under BRCE the overclaim set is empty as an invariant: an artifact without a receipt satisfying (B3) is not actuated. Hence the admissible magnitude of any claim a system makes is bounded by what its mechanism can receipt, not by its rhetoric.*

Proof. (B3) requires a receipt per actuation; an actuation without one violates BRCE and is, by (B1)’s gate, never emitted. So the set of receiptless actuations is empty, and every standing claim corresponds to a receipted artifact whose verification cost is $O(\kappa)$ (keystone). $\square$

This is the formal statement of the program’s discipline: *a grand system without receipts is grandiosity; a grand system that receipts its own limits is machinery.*

Chapter 7

A Worked Example: A Contract-Claim Law Object

We instantiate the whole apparatus on one object, tracing it $\text{Raw} \rightarrow \text{Receipted}$ with its obligations, an andon halt and its lift, and the resulting chain hash. The example is a *contract claim*: a payload asserting that a counterparty owes a delivery under a signed contract, submitted for lawful processing.

7.1 Raw

The caller submits an observation $o \in \mathcal{O}$, a JSON payload, to law judge:

```
{ "claim": "delivery-owed", "contract_id": "C-4471", "amount": 12000,
  "counterparty": "acme", "evidence": null }
```

The governing Law attaches three obligations (Definition 3.1), each a decidable g_i :

g_1 : `Precondition{predicate_id:"contract_active", params_hash: h(C-4471)}` — the referenced contract must be in the active set.

g_2 : `EvidenceRequired{evidence_type:"signed_contract"}` — a signed-contract record must be present.

g_3 : `BlockingConstraint{reason:"amount_within_authority"}` — the claimed amount must lie within the desk’s authority.

At submission the object is `LawObject<Claim,Raw,ContractLaw>` with `andon = Green` and `chain_hash = None`.

7.2 Judgment and an andon halt

`Judge::judge` evaluates g_1, g_2, g_3 . The contract C-4471 is active ($g_1 = 1$), but `evidence` is `null`: the signed-contract record is absent ($g_2 = 0$). By Definition 3.1, $\alpha(o) = \perp$. The arrow j returns the `Err` branch — the same object, in

`Andon::Halted{ unmet:[EvidenceRequired{"signed_contract"}], refusals:[MissingEvidence{...}], at`

By Proposition 3.3, $\text{cat}(\text{MissingEvidence}) = \text{Prerequisites}$ and $\text{lane}(\text{MissingEvidence}) = \text{PRECONDITION_FAILED}$; the total denial word is $d(o) = \text{PRECONDITION_FAILED} \neq \mathbf{0}$. This is a lawful output: a refusal with a category, a lane, and a timestamp — data, not an exception (Chapter 2). No manufacture occurs, because μ is undefined on $\perp$ (Definition 6.1); the line is halted, and on-red.

7.3 Evidence, re-judgment, admission

The caller supplies the missing evidence and resubmits:

```
{ ..., "evidence": { "signed_contract": "blob:9f2c..." } }
```

Now $g_2 = 1$ and $g_3 = 1$ (the amount 12000 is within authority), so $\bigwedge_i g_i = 1$ and $\alpha(o) = \rho(o) \in \mathcal{O}^*$: the object advances $j : \text{Raw} \rightarrow \text{Validated}$ to $\text{LawObject}\langle \text{Claim}, \text{Validated}, \text{ContractLaw} \rangle$ with $\text{andon} = \text{Green}$. Then $a : \text{Validated} \rightarrow \text{Admitted}$ admits it to **Admitted**. By Theorem 5.2 nothing could have jumped from **Raw** straight to **Admitted**; the two arrows were mandatory and in order.

7.4 Manufacture and receipt

The admitted claim is manufactured into its artifact (the actuated obligation-to-deliver record) and receipted. **receipt** consumes the **Admitted** object and appends the frame of Definition 6.2. With previous chain value h_- , a deterministic timestamp, instruction id, and the honest denial word **ADMITTED** (the halt was lifted, so the receipt records a clean admission), the payload is hashed $b \mapsto H(b)$ and the chain advances

$$h_+ = H(h_- \parallel \text{fr}(\theta, H(b))).$$

The object is now $\text{LawObject}\langle \text{Claim}, \text{Receipted}, \text{ContractLaw} \rangle$ carrying $\text{chain_hash} = \text{Some}(h_+)$. A lifecycle replay of $\text{judge} \rightarrow \text{admit} \rightarrow \text{receipt}$ yields fitness $0x0001.0000 = 1.0$ (Definition 6.3); all four BRCE invariants hold. The boundary projection the system exposes is the κ -sized tuple

$$r = (\text{verdict} = \text{pass}, h_+ = \langle \text{256-bit hex} \rangle, \varphi = 1.0, \text{reason} = \emptyset),$$

and by Theorem 6.3 this receipt is the unique chain position of this actuation, injective up to collision. A verifier trusts the delivery obligation by reading r — four chunks — never by re-deriving the manufacture.

Remark (The agent's byte). In the membrane deployment, each connected agent carries an eight-bit projection of its lifecycle status — the **agent8** byte — whose flags are set by exactly these events: a halt sets a **blocked** bit, a receipt sets a **receipted** bit. The worked object above drove one session's byte from **blocked** (at the andon halt) to **receipted** (after the chain advanced). The byte is the smallest receipt of all: a single octet, popcount-summable across a fleet, projecting a whole lifecycle into eight bits a scheduler can read in one instruction.

Chapter 8

Map of the Series

This foundations paper fixes the five objects, the Rice quarantine, the tpestate category, and BRCE. Four pillar papers develop each face in depth; the reader should approach them in any order after this one.

Pillar I — Admission Algebra. Deepens Chapter 2: the denial monoid D (Construction 3.1), the eight-category taxonomy as a total function (Proposition 3.3), and the algebra of obligation composition, including the join-irreducible structure of refusal categories and the prolog8 kernel’s proof-carrying admission path. It is the algebra beneath “refusal is data.”

Pillar II — Receipt Cryptography. Deepens Chapter 6, § 2: the BLAKE3 `OcelCausalFrame` chain (6.1), the Faithful Projection Theorem (a bounded verifier reads $O(1)$ symbols and rejects any perturbation except with negligible probability), the affidavit `verify` pipeline (`decode`, `check_format`, `chain_integrity`, `continuity`, `verify_commitments`, `evaluate_profile`), and receipt signing. It is the cryptography beneath Lemma 6.2 and Theorem 6.3.

Pillar III — Planning and Geometry. Develops manufacture as search and conformance as geometry: the `bcinr-pddl` planner that grounds and solves a manufactured PDDL8 domain, the marking polytope and state equation of token-flow execution, conformance as membership with Farkas separating-hyperplane certificates, and `PowlReplayVerifier` fitness (Definition 6.3) as normalized distance to the polytope. It is the geometry beneath the conformance invariant (B4).

Pillar IV — The Projection Principle (keystone). The already-standing keystone *The Projection Principle* unifies the above around the Comprehension–Verification Gap: verification cost is $O(\kappa)$ per receipt while comprehension cost is $\Omega(\dim \Sigma)$, so trust in a system of unbounded interior is available at constant human cost. BRCE (Definition 6.4) is that paper’s operating premise made into a checklist.

Prior work — [1]. The published paper *The Chatman Equation and the Industrial Revolution of Knowledge* (v1.2.0) demonstrated $\mathcal{A} = \mu(\mathcal{O})$ at enterprise scale: knowledge hooks $h =$ (trigger, check, act, receipt) over RDF/SHACL, the KNHK hot-path executor (≤ 2 ns rule

checks), the **ggen** bounded projection, Lockchain provenance, and full coverage of the forty-three workflow patterns [4]. This series factors that paper's μ into admission α and manufacture μ , names the receipt space $\mathcal{R}$ it produced, and supplies the first-principles proofs its scale rested on.

Chapter 9

Conclusion

We set out to give the series its floor. The Chatman equation $\mathcal{A} = \mu(\mathcal{O}^*)$ relates five named objects; admission is a computable retraction rather than a semantic decision, because Rice’s theorem forbids the latter (Theorem 2.1, Corollary 2.2); refusal is a first-class value carrying its category through a denial monoid whose taxonomy is a compiler-certified total function (Proposition 3.3); the lifecycle $\text{Raw} \rightarrow \text{Validated} \rightarrow \text{Admitted} \rightarrow \text{Receipted}$ is a free path category in which every illegal transition is an uninhabited hom-set, so “illegal-transition-is-a-type-error” is a theorem discharged before runtime (Theorem 5.2); manufacture is deterministic and bounded; the receipt is a collision-binding projection; and the enforced form BRCE yields conservation of consequence (Theorem 6.3) and its antibody corollary. A single contract-claim object exhibited all of it, halting on missing evidence and advancing to a receipted chain position once the evidence arrived.

The standard was never comprehension. It is a lawful admission, a bounded manufacture, and a faithful receipt — each small enough to hold, each guaranteed by mechanism. Everything downstream in this series is a deeper reading of one of those three.

Appendix A

Notation

Symbol	Meaning
$\mathcal{O}, \mathcal{O}^*$	observation space; admitted (decidable) sub-space $\mathcal{O}^*$
$\alpha, \perp$	admission retraction; refusal (first-class value)
$g_i, d(o)$	obligations; total denial word
$D = (\{0, 1\}^n, \vee, \mathbf{0})$	denial monoid (join-semilattice)
$\mu, \mathcal{A}$	manufacturing morphism; artifact/action space
$\Sigma, \mathcal{R}$	execution traces; receipt space
$H, h_{\pm}$	collision-resistant hash (BLAKE3); chain values
φ	conformance fitness $\in [0, 1]$ (Q16.16 in code; 1.0 = 0x0001_0000)
κ	human working-memory capacity (chunks), a fixed small constant
Life	lifecycle category on $\text{Raw} \rightarrow \text{Validated} \rightarrow \text{Admitted} \rightarrow \text{Receipted}$
j, a, r	judge, admit, receipt arrows
BRCE	Bounded Receipted Chatman Equation (B1–B4)

Appendix B

Code Correspondence

Every abstract object above is anchored to a location in the `praxis` repository.

Mathematical object	Code artifact	Location
$\mathcal{O}$, raw observation	<code>law judge</code> payload	<code>src/verbs/law.rs</code>
α , admission / $\perp$	<code>Judge/Admit</code> , <code>Andon</code>	<code>crates/praxis-core/src/law.rs</code>
g_i , obligations	Obligation enum (3 kinds)	<code>crates/praxis-core/src/law.rs</code>
D , denial monoid	<code>DenialPolarity</code> , <code>compose_denials</code>	<code>crates/praxis-core/src/refusal.rs</code>
cat, 8 categories	<code>RefusalCategory</code> , <code>::category</code>	<code>crates/praxis-core/src/refusal.rs</code>
Life , typestate	<code>LawObject<P,S:Stage,L></code> , sealed <code>Stage</code>	<code>crates/praxis-core/src/lifecycle.rs</code>
μ , manufacture	<code>mfg pddl</code> over <code>bcinr-pddl</code>	<code>src/verbs/mfg.rs</code>
h_+ , receipt chain	<code>OcelCausalFrame</code> , <code>build_admission_frame</code>	<code>crates/praxis-core/src/law.rs</code>
φ , fitness	<code>PowlReplayVerifier</code> , <code>replay_lifecycle</code>	<code>crates/praxis-core/src/replay_adapt.rs</code>
BRCE verify	<code>verify affidavit</code> pipeline	<code>crates/praxis-core/src/verify.rs</code>
<code>agent8</code> byte	resident <code>AgentByte</code> over MCP	<code>src/bin/mcp_lawobject_server.rs</code>

Bibliography

- [1] S. Chatman, *The Chatman Equation and the Industrial Revolution of Knowledge: $A = \mu(O)$, Knowledge Hooks, and Production-Verified Enterprise Execution*, v1.2.0, November 2025.
- [2] H. G. Rice, *Classes of recursively enumerable sets and their decision problems*, Transactions of the American Mathematical Society **74** (1953), 358–366.
- [3] R. E. Strom and S. Yemini, *Typestate: A programming language concept for enhancing software reliability*, IEEE Transactions on Software Engineering **SE-12**(1) (1986), 157–171.
- [4] W. M. P. van der Aalst, A. H. M. ter Hofstede, B. Kiepuszewski, and A. P. Barros, *Workflow patterns*, Distributed and Parallel Databases **14**(1) (2003), 5–51.
- [5] J. O’Connor, J.-P. Aumasson, S. Neves, and Z. Wilcox-O’Hearn, *BLAKE3: One function, fast everywhere*, 2020.
- [6] A. F. Ghahfarokhi, G. Park, A. Berti, and W. M. P. van der Aalst, *OCEL: A standard for object-centric event logs*, in New Trends in Database and Information Systems, Springer, 2021.
- [7] G. A. Miller, *The magical number seven, plus or minus two: some limits on our capacity for processing information*, Psychological Review **63**(2) (1956), 81–97.
- [8] N. Cowan, *The magical number 4 in short-term memory: A reconsideration of mental storage capacity*, Behavioral and Brain Sciences **24**(1) (2001), 87–114.