

Mission Geometry:
Planning, Conformance, and Separability
for Bounded Receipted Execution

The Chatman Equation Thesis Program
(Praxis / Capability Physics)

Genesis — Pillar III

Abstract

This is Pillar III of the Chatman Equation thesis series: the geometry beneath the manufacturing morphism μ and the conformance invariant of the enforced equation. We treat mission execution as two dual questions — *manufacture as search* (how a plan is produced) and *conformance as geometry* (how a trace is admitted) — and give both a first-principles account grounded line-for-line in the `praxis` and `bcinr-pddl` codebase. We define the PDDL action law with numeric fluents, and derive the attention fluent balance law as the feasibility condition of a durative-action schedule, with conservation of the borrowed-and-returned fluent proved directly. We build the marking polytope from the Petri state equation $M = M_0 + Nx$, show the STRIPS grounding is a place/transition net whose forward search inhabits marking space, and prove that the lifecycle token verifier’s branchless firing rule $\text{enabled} \leftarrow (\text{enabled} \wedge \neg \mathbf{m}^-) \vee \mathbf{m}^+$ is exactly the safe-net specialization of that equation. We formalize conformance as membership in the reachability cone, and prove a Farkas separation theorem: a nonconformant marking admits a separating-hyperplane certificate of dimension p , sound for non-reachability and verifiable at $O(\kappa)$. We define token-replay fitness as the population ratio the `PowlReplayVerifier` computes in Q16.16, and prove that unit fitness with completion characterizes exactly the genuine firing sequences, with the first disabled transition a coordinate-localized witness. We formalize POWL 2.0 as recursive partial orders and choice graphs, show the plan-to-tape transitive reduction yields the Hasse diagram of the precedence order, and reframe the Kourani–van der Aalst separable decomposition [?] as a Rice quarantine for processes: separable is admitted, non-separable is refused with reason. Finally we derive lexicographic cost arbitration as the infinite-weight limit in which the admitted coordinate dominates, and give the makespan functional with its critical-path and capacity-sensitivity structure as computed by `analyze_schedule`. Every theorem is anchored to the `bcinr-pddl` planner (`find_plan`, `find_temporal_plan`), the `PowlReplayVerifier`, the capability router `CostVector`, and the `ScheduleAnalysis64` analyzer.

Contents

Chapter 1

Introduction: Manufacture as Search, Conformance as Geometry

1.1 Where this pillar sits

The foundations paper of this series fixed the Chatman equation $\mathcal{A} = \mu(\mathcal{O}^*)$ and its enforced form, the Bounded Receipted Chatman Equation (BRCE there), a conjunction of four invariants: an admission gate, a bounded deterministic manufacture, receipt totality, and *conformance* — the requirement that an actuation’s lifecycle replay have fitness $\varphi = 1$. Two of those invariants are the subject of this pillar. The manufacturing morphism μ , when its artifact is a *mission* — a schedule of durative actions achieving a goal — is a bounded search over a state space. The conformance invariant is a *geometric* membership test: does a trace lie inside the polytope of lawful markings? We develop both from first principles and show they are the same object viewed from two sides.

Question	Object	praxis/bcinr-pddl realization
Manufacture (μ)	bounded search	<code>find_plan</code> (BFS), <code>find_temporal_plan</code>
Conformance (B4)	cone membership	<code>PowlReplayVerifier</code> , marking geometry
Arbitration	lexicographic order	<code>CostVector</code> , <code>route_capability_plan</code>
Makespan / cost	critical-path functional	<code>ScheduleAnalysis64</code>

1.2 The two theses of this pillar, stated once

Manufacture is bounded search over a finite ground net. A lifted PDDL domain with numeric fluents grounds to a *finite* set of transitions (Theorem ??), so μ ’s search terminates with a priori bounded cost — the boundedness law (M2) of the foundations, instantiated. The state it searches is a marking: a vector of atom-tokens plus a valuation of numeric fluents. The attention fluent, borrowed at an action’s start and returned at its end, obeys a balance law whose nonnegativity is exactly schedule feasibility (Proposition ??).

Conformance is distance to a cone, certified by a hyperplane. The reachable markings of a net are constrained by the state equation $M = M_0 + Nx$ to lie in a translated cone (Proposition ??). A trace conforms only if its marking lies there; a nonconformant marking is

separated from the cone by a Farkas hyperplane (Theorem ??), a certificate of dimension p that a bounded verifier checks without replaying the interior. Token-replay fitness measures the normalized excursion outside the enabled set, and equals 1 exactly on genuine firing sequences (Theorem ??).

Throughout, a claim is a definition, a theorem with proof, or a citation. The prior published paper [?] demonstrated $\mathcal{A} = \mu(\mathcal{O})$ at enterprise scale and full coverage of the forty-three workflow patterns [?]; this pillar supplies the planning and conformance geometry those numbers rested on.

Chapter 2

The PDDL Action Law with Numeric Fluents

2.1 Lifted domains, grounding, and the state

We take the numeric-temporal fragment of PDDL 2.1 [?] as the language of μ when it manufactures missions, and give it exactly the structure `bcinr-pddl` grounds and searches.

Definition 2.1 (Lifted numeric-temporal domain). A *lifted domain* is a tuple $\mathcal{D} = (\mathcal{T}, \mathcal{P}, \mathcal{F}, \mathcal{S})$: a finite type hierarchy $\mathcal{T}$ (a forest under a subtype relation, with universal root `object`); a finite set $\mathcal{P}$ of predicate symbols with typed arities; a finite set $\mathcal{F}$ of numeric *function* symbols (fluents) with typed arities; and a finite set $\mathcal{S}$ of action schemas. A *durative* schema $s \in \mathcal{S}$ carries typed parameters $\bar{v}$, a duration constraint, a condition C_s , and an effect E_s , where conditions and effects are timestamped `at start` / `at end` and may be atomic (over $\mathcal{P}$) or numeric comparisons/updates (over $\mathcal{F}$).

Definition 2.2 (Grounded temporal problem). Given a domain $\mathcal{D}$ and a problem $(\mathcal{O}, M_0, \nu_0, \gamma)$ with finite typed objects $\mathcal{O}$, initial atom set M_0 , initial fluent valuation $\nu_0 : \mathcal{F}\text{-instances} \rightarrow \mathbb{R}$, and goal condition γ , the *grounding* substitutes every type-compatible object tuple into each schema’s parameters, producing a finite set of ground actions. A grounded state is a pair (M, ν) of a set M of true ground atoms and a fluent valuation ν .

This is the `GroundTemporalProblem` of `bcinr-pddl`: `initial_atoms`, `initial_fn_values`, `grounded_durative_actions`, and a goal `PddlCondition`. Type compatibility is the `TypeIndex::satisfies` walk up the parent chain; the universal type `object` matches every object.

Theorem 2.1 (Grounding is bounded). *Let $\mathcal{D}$ have schemas of maximum parameter arity k over $|\mathcal{O}| = N$ objects. The number of ground actions is at most $|\mathcal{S}| \cdot N^k$, a finite constant independent of the goal. Consequently a forward search over grounded actions has a branching factor bounded a priori, and μ ’s manufacture terminates within a fixed depth bound.*

Proof. Each schema binds at most k parameters, each to one of $\leq N$ objects, so it grounds to at most N^k actions; summing over $|\mathcal{S}|$ schemas gives the bound. In `bcinr-pddl` the count is additionally capped by the structural constant `PDDL8_MAX_GROUND`: exceeding it is the hard error `Pddl8Error::BoundExceeded`, and an empty grounding is `EmptyGrounding`. Thus the ground action set is finite and its size is a fixed function of $(|\mathcal{S}|, N, k)$, discharging the boundedness law (M2) for mission manufacture. $\square$

Remark (The forward search is the manufacture). For the classical (atomic) fragment, `GroundProblem::find_plan` is breadth-first search over sets of ground atoms: the frontier is a queue of (state, path) pairs, visited states are deduplicated, and the first state containing the goal set returns its path as a `Pddl8Tape`. Search failure within `PDDL8_MAX_PLAN_DEPTH` is `NoAdmittedPlan` — a *refusal with reason*, not an exception, exactly as admission returns $\perp$. The manufacture either yields an artifact or a receipted refusal; it never diverges, by Theorem ??.

2.1.1 Relational Grounding and XOR Filter Membership Gates

To avoid the naive grounding parameter explosion, grounding is refactored into a relational database join query over type-safe dictionary-encoded symbol IDs: $\text{SymId} \in \mathbb{Z}_{>0}$.

Construction 2.1 (XOR-Filter-Pruned Membership Gate). Let $R \subseteq \mathcal{O}$ be the reachability fixpoint fact set. A frozen fact store builds an 8-bit XOR filter over FNV-1a hashes of R . Index mapping uses the fastrange reduction:

$$\text{reduce}(x, n) = \frac{x \times n}{2^{32}}, \quad (2.1)$$

which projects hash x onto n blocks using multiplication and bit-shifts. Let F be the filter and B the sorted BTreeSet of R . Membership is gated by:

$$\text{contains}(x) = \begin{cases} x \in B, & \text{if } F(x) = \text{true}, \\ \text{false}, & \text{otherwise.} \end{cases}$$

The false-positive rate is $\approx 0.4\%$, and false negatives are 0.

2.2 Numeric fluents and the attention balance law

The distinctive content of the numeric fragment is that an action may test and update a real-valued fluent. `bcinr-pddl`'s `eval_numeric` evaluates a fluent expression against a valuation, and `apply_numeric_effect` realizes the five update kinds `Assign`, `Increase`, `Decrease`, `ScaleUp`, `ScaleDown`. We formalize the one that carries the calculus of attention.

Definition 2.3 (Numeric fluent balance). A durative action j *draws* on a fluent ϕ if its effect contains **at start** (`decrease  $\phi$   $r_j$`) and **at end** (`increase  $\phi$   $r_j$`) for a rate $r_j \geq 0$, and its condition contains **at start** (`$\phi \geq r_j$`). Over the half-open interval $[s_j, e_j)$ of its execution, j holds r_j units of ϕ and releases them at e_j . The *free level* of ϕ at time t is

$$f_\phi(t) = \nu_0(\phi) - \sum_{j: t \in [s_j, e_j)} r_j.$$

This is exactly the capability router's model: its domain declares `(:functions (attention));` each capability schema decrements `attention` by 1 **at start** under the guard `(>= (attention) 1)` and increments it back **at end**. Attention is “one working human, one active train of thought.”

Proposition 2.2 (Attention is borrowed and returned: conservation). *If every action drawing on ϕ both decrements ϕ by r_j at start and increments it by r_j at end, then the net effect of each completed action on ϕ is zero, and the terminal valuation equals the initial: $\nu_{\text{end}}(\phi) = \nu_0(\phi)$. Hence ϕ is a conserved resource; scheduling redistributes when it is held, never how much exists.*

Proof. By `apply_numeric_effect`, `Decrease`(ϕ, r_j) then `Increase`(ϕ, r_j) compose to the identity on $\nu(\phi)$. Summing over all completed actions, the total change is $\sum_j (r_j - r_j) = 0$, so $\nu_{\text{end}}(\phi) = \nu_0(\phi)$. The intermediate values differ only on the union of the open execution intervals, which is the content of Definition ??.

Proposition 2.3 (Feasibility is pointwise nonnegativity). *A durative-action schedule is feasible with respect to ϕ — every action’s start guard ($\phi \geq r_j$) holds when it fires — if and only if $f_\phi(t) \geq 0$ for all t ; equivalently, at no instant does the total rate of in-flight draws exceed $\nu_0(\phi)$. For the unit-rate attention fluent ($r_j \equiv 1$), this says the number of concurrently in-flight capabilities never exceeds the capacity $\nu_0(\text{attention})$.*

Proof. The start guard at s_j requires $\nu(\phi) \geq r_j$ in the state just before j ’s at-start decrement. By Definition ??, that state’s value of ϕ is $f_\phi(s_j^-) = \nu_0(\phi) - \sum_{i \neq j: s_j \in [s_i, e_i)} r_i$, so the guard holds iff $f_\phi(s_j^-) \geq r_j$, i.e. $f_\phi(s_j) \geq 0$ after j ’s own draw is included. Free level changes only at action starts and ends, so if it is ≥ 0 at every start it is ≥ 0 everywhere. Conversely a violated guard is an instant with $f_\phi < 0$. For $r_j \equiv 1$, the sum counts in-flight actions, so nonnegativity is exactly “concurrency $\leq$ capacity.”

`bcinr-pddl`’s temporal planner enforces exactly this: within a tick it re-scans candidate actions after each start so an `at start` decrement is visible to the next candidate’s guard in the same instant, which the source comments identify as “what lets concurrent starts correctly gate on shared resource fluents.” The `zero_attention_capacity_is_infeasible_and_refused` test is Proposition ?? at capacity 0: with $f_{\text{attention}}(t) < 0$ forced at any start, the router refuses rather than defaults a route.

Chapter 3

The Marking Polytope and the State Equation

3.1 Place/transition nets and the incidence matrix

Atomic (non-numeric) execution has a native linear-algebraic geometry, that of place/transition nets [?]. We reproduce it in the form the code inhabits.

Definition 3.1 (Net, marking, enabling, firing). A *net* has p places and a finite set T of transitions. A *marking* is a vector $M \in \mathbb{Z}_{\geq 0}^p$. A transition t is a pair $(\mathbf{m}_t^-, \mathbf{m}_t^+)$ of a *preset* (consumed) and *postset* (produced) vector. t is *enabled* at M iff $M \geq \mathbf{m}_t^-$ coordinatewise; firing it yields $M' = M - \mathbf{m}_t^- + \mathbf{m}_t^+$.

Definition 3.2 (Incidence matrix and state equation). Let $N \in \mathbb{Z}^{p \times |T|}$ have column $\delta_t = \mathbf{m}_t^+ - \mathbf{m}_t^-$ for each transition t . A firing sequence ς with *Parikh vector* $x \in \mathbb{Z}_{\geq 0}^{|T|}$ (the count of each transition's firings) satisfies the *state equation*

$$M = M_0 + Nx. \quad (3.1)$$

Proposition 3.1 (Reachability lies in a translated cone). *Every marking M reachable from M_0 satisfies (??) for some $x \in \mathbb{Z}_{\geq 0}^{|T|}$, hence lies in the integer points of $M_0 + \text{cone}(N)$ intersected with $\mathbb{Z}_{\geq 0}^p$, where $\text{cone}(N) = \{Nx : x \geq \mathbf{0}\}$. The state equation is a necessary condition for reachability; it is not sufficient (a x solving (??) need not correspond to a fireable ordering).*

Proof. Firing t once adds column δ_t to the marking (Definition ??); a sequence firing each t exactly x_t times adds $\sum_t x_t \delta_t = Nx$, giving (??). Since each $x_t \geq 0$ and markings are nonnegative integers, M is an integer point of $M_0 + \text{cone}(N)$ in the nonnegative orthant. Insufficiency is the classical Petri-net fact [?]: (??) ignores intermediate nonnegativity and firing order, so spurious solutions exist. $\square$

Definition 3.3 (Marking polytope). The *marking polytope* of a net is the relaxation

$$\mathcal{P} = \{ M_0 + Nx : x \geq \mathbf{0} \} \cap \{ M \geq \mathbf{0} \} \subseteq \mathbb{R}_{\geq 0}^p,$$

the continuous translated cone whose integer points over-approximate the reachable set.

3.2 STRIPS grounding is a net; the lifecycle token model is its safe specialization

Construction 3.1 (STRIPS grounding as a net). Take the places to be the ground atoms. A ground action t with delete-effects $\mathbf{m}_t^-$ and add-effects $\mathbf{m}_t^+$ is a transition; its precondition atoms must all be present for it to fire (an enabling condition refining Definition ??). Then `GroundProblem::find_plan` searches marking space: the BFS frontier is a set of markings, the successor relation is exactly transition firing (`for d in del_effects { next.remove(d) }`; `for a in add_effects { next.insert(a) }`), and a plan is a path to a marking containing the goal.

The lifecycle model of the foundations paper — the fixed sequence `judge` $\rightarrow$ `admit` $\rightarrow$ `receipt` — is a concrete net whose places are the token bits of `replay_adapter`: `TOK_START`, `TOK_JUDGED`, `TOK_ADMITTED`, `TOK_DONE`. Each step's `required_tokens` is its preset and `produces_tokens` its postset:

`judge` : `TOK_START` $\rightarrow$ `TOK_JUDGED`, `admit` : `TOK_JUDGED` $\rightarrow$ `TOK_ADMITTED`, `receipt` : `TOK_ADMITTED` $\rightarrow$ `TOK_DONE`

Proposition 3.2 (The verifier's firing rule is the safe-net state equation). *On a safe (1-bounded) net, where every marking and every preset/postset is a 0/1 vector and presets are consumed, the integer firing rule $M' = M - \mathbf{m}_t^- + \mathbf{m}_t^+$ of Definition ?? coincides with the branchless bitset update the `PowlReplayVerifier` performs:*

$$enabled_tokens \leftarrow (enabled_tokens \& \neg \mathbf{m}_t^-) \mid \mathbf{m}_t^+,$$

and the enabling test $M \geq \mathbf{m}_t^-$ coincides with the branchless subset check $(enabled \& \mathbf{m}_t^-) \oplus \mathbf{m}_t^- = 0$.

Proof. On 0/1 markings with $\mathbf{m}_t^- \leq M$, the AND-NOT `enabled` $\& \neg \mathbf{m}_t^-$ clears exactly the consumed places (subtracting $\mathbf{m}_t^-$), and the OR with $\mathbf{m}_t^+$ sets the produced places (adding $\mathbf{m}_t^+$); when preset and postset are disjoint from the retained marking this is bit-for-bit the integer expression. The enabling identity holds because $(M \& \mathbf{m}_t^-) \oplus \mathbf{m}_t^- = 0$ iff every bit of $\mathbf{m}_t^-$ is also set in M , i.e. $M \geq \mathbf{m}_t^-$ coordinatewise on 0/1 vectors. Both are the verifier's two guard lines verbatim. $\square$

Proposition ?? is the bridge between the general Petri geometry of [?] and the constant-time bitwise mechanism the receipt path actually runs: the polytope theory is the semantics, the mask reduction is the implementation.

Chapter 4

Conformance as Cone Membership and Farkas Certificates

4.1 Conformance is membership

Definition 4.1 (Conformance). A trace with final marking M_f and transition-count vector x *conforms* to a net with initial marking M_0 if $M_f \in \mathcal{P}$ and the x witnessing it is realized by a genuine firing ordering. A trace *fails membership* if $M_f \notin \mathcal{P}$, i.e. no $x \geq \mathbf{0}$ solves $Nx = M_f - M_0$ within the nonnegative orthant.

Theorem 4.1 (Farkas separation certifies nonconformance). *Fix $b = M_f - M_0$. Exactly one of the following holds:*

- (a) *there exists $x \geq \mathbf{0}$ with $Nx = b$ (the continuous relaxation is feasible; membership in the cone is possible); or*
- (b) *there exists a vector $y \in \mathbb{R}^p$ with $N^\top y \leq \mathbf{0}$ and $y^\top b > 0$ — a separating hyperplane whose normal y certifies $M_f \notin M_0 + \text{cone}(N)$.*

The certificate y is a sound proof of nonconformance: if (b) holds, no firing sequence, in any order, reaches M_f from M_0 .

Proof. This is Farkas' lemma [?] applied to the system $Nx = b$, $x \geq \mathbf{0}$: either the system has a nonnegative solution, or there is y with $y^\top N \leq \mathbf{0}^\top$ and $y^\top b > 0$, and never both. In case (b), suppose for contradiction a firing sequence reached M_f ; by Proposition ?? its Parikh vector $x^* \geq \mathbf{0}$ satisfies $Nx^* = b$, so $y^\top b = y^\top Nx^* = (N^\top y)^\top x^* \leq 0$, contradicting $y^\top b > 0$. Hence (b) rules out reachability. Soundness is one-directional by Proposition ??: feasibility of the relaxation (a) is necessary but not sufficient for an actual firing sequence, so (a) does not by itself certify conformance — only (b) certifies its negation. $\square$

Corollary 4.2 (The certificate is bounded and verifiable at $O(\kappa)$). *A nonconformance certificate is a single vector $y \in \mathbb{R}^p$ together with the two inequalities $N^\top y \leq \mathbf{0}$ and $y^\top b > 0$. Checking it is a matrix product and two sign tests — independent of the length of the trace that produced M_f . Projected to the receipt's bounded coordinate (a verdict and a localized witness), it is verifiable within a bounded context κ , per the keystone's Comprehension–Verification Gap.*

Proof. Verification reads y and b and evaluates $N^\top y$ and $y^\top b$; the cost is $O(p \cdot |T|)$ in the net's fixed dimensions, with no dependence on trace length or interior computation. The receipt exposes the Boolean verdict and the witness coordinate, of size $\leq \kappa$. $\square$

Remark (Geometry is the process-layer projection principle). A conformance failure is not reported as “the trace looked wrong”; it is a hyperplane one can recompute. This is the projection principle in linear algebra: an unbounded execution is judged by a bounded certificate whose fidelity is mechanical, not a matter of the verifier comprehending the run.

Chapter 5

Token-Replay Fitness as Normalized Distance

5.1 The replay verifier as a Petri semantics

Token-replay conformance [?] plays a trace against a process model and measures how far it strays from lawful firing. `bcinr-powl`'s `PowlReplayVerifier` is a branchless realization; we give its exact semantics and prove what it characterizes.

Definition 5.1 (Replay state and step). The verifier holds a marking M (`enabled_tokens`) initialized to the entry token, and accumulators `replayed`, `fitted`, `enabled_not_taken` (bitsets of node identities). A frame carries a one-hot `node_bit`, a preset `required_tokens` = $\mathbf{m}^-$, and a postset `produces_tokens` = $\mathbf{m}^+$. Replaying a frame:

- (1) rejects a `node_bit` that is zero or not a power of two (`UnknownNode`);
- (2) requires $M \geq \mathbf{m}^-$ (else `TokenNotEnabled`), the branchless check of Proposition ??;
- (3) records the enabled-but-unconsumed tokens `enabled_not_taken` $\mid= M \& \neg \mathbf{m}^- \& \neg \text{node_bit}$;
- (4) fires: $M \leftarrow (M \& \neg \mathbf{m}^-) \mid \mathbf{m}^+$; and
- (5) sets the `node_bit` in both `replayed` and `fitted`.

Definition 5.2 (Conformance fitness). On finalization the verifier returns, in Q16.16 fixed point,

$$\varphi = \frac{|\text{fitted}|}{|\text{replayed}|}, \quad \text{precision} = \frac{|\text{replayed}|}{|\text{replayed}| + |\text{enabled_not_taken}|},$$

where $|\cdot|$ is population count and $\varphi := 1$ when the denominator is zero. The value 1.0 is the constant `0x0001_0000`.

Theorem 5.1 (Unit fitness characterizes genuine firing sequences; the fault localizes). *Replaying a sequence of frames against the net either (i) completes with no violation, in which case the sequence is a genuine firing sequence — each transition was enabled at the marking where it fired — and $\varphi = 1$; or (ii) halts at the first frame whose preset is not contained in the current marking, returning `TokenNotEnabled`{`node_id`}, a coordinate-localized witness of the earliest disabled transition.*

Proof. Step (2) of Definition ?? admits a frame iff $M \geq \mathbf{m}^-$, which by Definition ?? is exactly the enabling condition; step (4) is the firing rule (Proposition ??). By induction on the prefix, if no frame is rejected then every frame fired from a marking at which it was enabled, so the whole sequence is a genuine firing sequence and the intermediate markings are reachable. In that case every replayed frame also sets **fitted** (step 5), so $|\mathbf{fitted}| = |\mathbf{replayed}|$ and $\varphi = 1$. Conversely the guard in step (2) fails at the *first* frame whose required tokens are absent — the verifier returns immediately with that frame’s **node_id** — so a non-firing sequence is rejected at the earliest offending index, which is the localized witness t^* . $\square$

Proposition 5.2 (Fitness and precision are honest population ratios). *φ and precision are ratios of bit-populations of disjointly-maintained bitsets; neither can exceed 1, and $\varphi = 1$ is attained exactly on completed replays (Theorem ??). The **enabled_not_taken** accumulator, gathered before consumption in step (3), counts tokens that were live but not used by the firing — the model’s under-fitting — and feeds precision, not fitness; the two metrics therefore measure orthogonal deviations and cannot be traded against each other.*

Proof. Population count is monotone and bounded by the number of node bits, and **fitted** $\subseteq$ **replayed** by construction (step 5 sets both), so $\varphi \leq 1$ with equality iff no frame was replayed-without-fitting — which, in the current step, is every completed replay. **enabled_not_taken** is updated only in step (3) and read only by precision’s denominator, so it is independent of the fitness numerator. $\square$

The lifecycle instance makes this concrete. **replay_lifecycle** on the in-order sequence [Judged, Admitted, Receipted] returns $\varphi = 0x0001_0000$; the out-of-order [Admitted, ...] fails at the Admitted frame, which requires TOK_JUDGED that the entry marking lacks, returning **TokenNotEnabled**{*node_id* : 2}; and skipping **admit** before **receipt** fails at node 3. These are the three tests **lawful_in_order_sequence_has_fitness_one**, **out_of_order_sequence_is_token_not_enabled**, and **skipping_admit_before_receipt_is_token_not_enabled** — Theorem ??’s two branches, executed.

Remark (Relation to the keystone’s fractional fitness). The keystone defines $\varphi = 1 - (\text{tokens forced on unenabled tr})$ a fraction in $[0, 1]$. The **praxis** lifecycle verifier is the strict specialization in which a forced-disabled firing is a hard **TokenNotEnabled** rather than a fractional penalty: the “forced consumption” branch is never taken because the verifier refuses to fire a disabled transition at all. Both characterize the same lawful set $\{\varphi = 1\}$; the specialization is more conservative — it localizes the first fault and stops — which is the desired behavior for a lifecycle whose only lawful order is a single line.

Chapter 6

POWL 2.0: Choice Graphs and Partial Orders

6.1 The language

Partially Ordered Workflow Language (POWL) [?] is the model class into which the planner's output is compiled for replay. We define it as the recursive structure the code builds.

Definition 6.1 (POWL model). A *POWL model* over an activity alphabet A is one of:

- (i) an *activity* $a \in A$ (a leaf);
- (ii) a *partial order* $(\{M_1, \dots, M_k\}, \prec)$: a finite set of child POWL models with an irreflexive, acyclic precedence relation $\prec$, executed by respecting $\prec$ and running $\prec$ -incomparable children concurrently; or
- (iii) a *choice graph* $(\{M_1, \dots, M_k\}, E)$: children combined by exclusive choice (and, in POWL 2.0, cyclic/loop edges E), of which the execution takes exactly one live branch.

The three node kinds are `PowlOpKind::Activity`, `PartialOrderGate`, `ChoiceGate` in `powl_bridge`.

Construction 6.1 (Plan to POWL tape; transitive reduction). `temporal_plan_to_powl_tape` converts a `TemporalPlan` into a tape of `PowlOpSpecs`. Each step i becomes an activity with a `pred_mask` initialized to the set of steps $j < i$ that finish before i starts (`prev_end ≤ starti`), and a one-hot `succ_mask = 1 ≪ i`. A second pass performs *transitive reduction*: for each predecessor k of i , the bits of k 's own predecessors are cleared from i 's mask, leaving only *direct* predecessors.

Proposition 6.1 (The reduced masks are the Hasse diagram of the precedence order). *Let $\prec$ be the strict order “ j finishes before i starts” on plan steps. The initialized `pred_mask` of step i is $\{j : j \prec i\}$, and after transitive reduction it is the covering relation of $\prec$: j remains a predecessor of i iff $j \prec i$ and there is no k with $j \prec k \prec i$. The resulting dependency graph is a DAG, and two steps are schedulable concurrently iff they are $\prec$ -incomparable.*

Proof. $\prec$ is irreflexive and transitive (it is induced by the total order on real finish/start times), hence a strict partial order. The first pass sets bit j in i 's mask exactly when $j \prec i$. The reduction clears bit j from i 's mask when some retained predecessor k of i has j in its

predecessors, i.e. $j \prec k \prec i$; what survives is precisely the covering relation (the Hasse diagram) of $\prec$. Acyclicity follows from irreflexivity and transitivity. Incomparable steps have no path between them, so the scheduler may overlap their intervals; comparable steps are ordered by $\prec$. $\square$

Proposition 6.2 (Local marking dimension is bounded by node arity). *In a POWL model, the token marking local to a node ranges over the tokens of that node's immediate children: a partial-order or choice node of arity k has a local marking space of dimension $O(k)$, independent of the total size of the model. Consequently the replay of a node is checked in a working set bounded by its arity, and the global replay is the composition of these bounded-dimension checks.*

Proof. By Definition ?? a node's execution semantics reference only its own children's tokens: a partial order gates on the $\prec$ -predecessors among its k children; a choice gates on which of its k branches is live. The token bits touched by one node's firing are therefore contained in an $O(k)$ -dimensional subspace of the marking, regardless of how large the surrounding model is. The bounded-arity constant is the `PowlOpSpec` tape's per-op mask width. $\square$

Proposition ?? is the structural reason conformance is cheap: even though the mission's interior may be arbitrarily large, each POWL node is judged in a working set bounded by its arity — the projection principle, localized to the process tree.

Chapter 7

Separability as a Rice Quarantine for Processes

7.1 The separable decomposition

Not every workflow net is a POWL model; the ones that are, are exactly the *separable* ones, and this is a decidable structural property.

Theorem 7.1 (Separable decomposition, after Kourani–van der Aalst). *A sound, separable workflow net admits a recursive decomposition into a POWL model (Definition ??) whose every internal node is a partial order (all concurrent / sequential logic) or a choice graph (all exclusive / cyclic logic), and whose leaves are activities. Each node’s local marking geometry has dimension bounded by its arity (Proposition ??), independent of the global net size.*

Proof. We prove the existence of the recursive decomposition by induction on the number of transitions $|T|$ of the workflow net W . For the base case, if $|T| = 1$, the workflow net consists of a single transition t connected to the source place and the sink place. This net is decomposed into a leaf POWL model corresponding to the activity $a = \text{label}(t)$, satisfying the theorem statement. For the inductive step, assume the claim holds for all sound, separable workflow nets with fewer than n transitions. Let $W = (P, T, F)$ be a sound, separable workflow net with $|T| = n > 1$. Since W is separable, it contains a top-level separator (a place, a transition, or a boundary set) that partitions the transitions T into k disjoint subsets $T_1, \dots, T_k$ such that the sub-nets induced by these partitions are independent and interact only at the separator boundary. By the definition of separability, this partition corresponds to one of the canonical operators: a sequential execution (partial order with total order), a concurrent execution (partial order with incomparable components), or an exclusive choice (choice graph). For each partition T_i , we construct a normalized workflow sub-net W_i using the Sub-Net Normalization Interface (Construction ??), which adds fresh start and end boundary places. Since the original workflow net W is safe and sound, each normalized sub-net W_i is also safe and sound, and satisfies $|T_i| < n$. By the induction hypothesis, each sub-net W_i admits a recursive decomposition into a POWL model ψ_i . We then define the POWL model for W as the composition of the sub-models: $\psi = \text{OP}(\psi_1, \dots, \psi_k)$, where OP is a partial order operator or a choice graph operator matching the separator type. For the local marking geometry: since the children in the decomposition are disjoint sub-nets, the token flow at the parent node is restricted to the boundary places and transitions of the children. The dimension of this local marking space is bounded by $O(k)$,

where k is the arity of the operator, and does not depend on the inner structure of the children or the global net size. This completes the proof. $\square$

Proposition 7.2 (Separability is the process-layer Rice quarantine). *Let $\text{sep} : \{\text{workflow nets}\} \rightarrow \{0, 1\}$ be the (decidable) predicate “the net is sound and separable.” Then sep is a legitimate admission obligation g in the sense of the foundations paper: it is total and computable, so admitting a process by sep retracts the space of arbitrary control-flow onto the decidable sub-language of POWL-expressible processes. A separable net is admitted (it decomposes, and its conformance is the bounded-arity replay of Chapter ??); a non-separable net is refused with reason — the reason being the structural obstruction to decomposition.*

Proof. Soundness and separability are structural properties of a finite net, checkable by the decomposition procedure of Theorem ??, which either returns a POWL tree or the node at which decomposition fails. Thus sep is total and computable, satisfying the obligation criterion (a syntactic terminating check, never a semantic decision about what the process “means”). Admission maps the net to its POWL decomposition when $\text{sep} = 1$ and to $\perp$ with the obstruction as refusal data otherwise, which is exactly the admission map α specialized to processes. $\square$

Remark (The parallel is exact). The foundations paper retracted the undecidable observation space $\mathcal{O}$ onto a decidable admitted language $\mathcal{O}^*$ because Rice’s theorem forbids deciding meaning. Here an unwieldy space of arbitrary control-flow is retracted onto the decidable sub-language of separable (POWL) processes because only there is conformance a bounded-arity replay. In both cases: a large, badly-behaved space is replaced by a small decidable sub-language whose membership is mechanically certified, and the boundary certificate — a POWL tree, a fitness value, a Farkas hyperplane — is small. Separability *is* admission, for processes.

Construction 7.1 (Sub-Net Normalization Interface). Given a partition part $T' \subseteq T$ with entry places P_{in} and exit places P_{out} , the sub-net is normalized to a valid WF-net by adding fresh boundary places p_s, p_e and silent transitions $\tau_{\text{in}}, \tau_{\text{out}}$. The normalized flow relation is:

$$F_{\text{norm}} = F' \cup \{(\text{src}, \tau_{\text{in}}), (\tau_{\text{in}}, p_s), (p_e, \tau_{\text{out}}), (\tau_{\text{out}}, \text{snk})\}, \quad (7.1)$$

redirecting boundary arcs onto the unified source src and sink snk .

Theorem 7.3 (Recursive Language Correctness). *Let W be a safe $\mathcal{E}$ sound workflow net, and let $\psi = \text{convert}(W)$ be its POWL 2.0 conversion. Then for any prefix trace length k :*

$$\mathcal{L}_k(W) = \mathcal{L}_k(\psi) = \mathcal{L}_k(\text{recompose}(\psi)).$$

Proof. We establish the language equivalence $\mathcal{L}_k(W) = \mathcal{L}_k(\psi)$ by induction on the depth of the recursive POWL structure ψ . For the base case of depth zero, ψ is a single activity leaf representing a transition t in W . The language $\mathcal{L}_k(W)$ of the net containing only t is the set of trace prefixes of length $\leq k$ over the alphabet $\{t\}$, which is exactly the trace language of the leaf activity $\mathcal{L}_k(\psi)$. For the inductive step, assume the equivalence holds for all child POWL sub-models of depth less than d . Let $\psi = \text{OP}(\psi_1, \dots, \psi_m)$ be a POWL model of depth d , and let $W_1, \dots, W_m$ be the normalized workflow sub-nets of W corresponding to the children. We analyze based on the operator type of OP : If OP is a partial order operator with precedence relation $\prec$, the language $\mathcal{L}_k(\psi)$ is the set of all prefixes of length $\leq k$ of sequences in the shuffle

of the languages of the children ψ_i that respect $\prec$. By the separability of W , the transition partitions T_i are disjoint and share no places except the boundary interface. Thus, any firing sequence in W interleaves transitions of concurrent sub-nets without local state conflicts, while enforcing the precedence $\prec$ through the boundary place connections, which yields exactly $\mathcal{L}_k(W) = \mathcal{L}_k(\psi)$. If OP is a choice graph operator, the language $\mathcal{L}_k(\psi)$ is the union of the prefix trace languages of the selected branches. In the net W , the entry place of the choice component is safe and holds a single token, which enforces that at most one branch sub-net can be activated and fire, preventing concurrent execution of different branches. Furthermore, cycle edges in the choice graph correspond to returning a token to the entry place, which matches the token game cyclic replay. Therefore, the language of the choice net matches the choice graph trace language, yielding $\mathcal{L}_k(W) = \mathcal{L}_k(\psi)$. Finally, by Construction ??, the recomposition operator ‘recompose’ reconstructs the net structure by recursively replacing POWL nodes with their canonical Place/Transition net templates. Since the templates preserve the exact flow relation and boundary places, the resulting net $\text{recompose}(\psi)$ is isomorphic to W , which guarantees $\mathcal{L}_k(W) = \mathcal{L}_k(\text{recompose}(\psi))$. Thus, the languages are identical for all k . $\square$

Chapter 8

Cost Arbitration and the Makespan Functional

8.1 The cost vector and lexicographic order

When several admitted plans achieve a goal, the router must choose one deterministically. The choice is a total order on a cost vector.

Definition 8.1 (Cost vector). The router's `CostVector` is the tuple

$$c = (c_0, c_1, c_2, c_3, c_4, c_5) = (\overline{\text{admitted}}, \text{risk}, \text{attention_seconds}, \text{tokens}, \text{latency}, \text{switches}),$$

where $c_0 = \overline{\text{admitted}}$ is the unadmitted indicator ($0 = \text{admitted}$, and `admitted` sorts first), $c_1 = \text{unreceipted_mutation_risk} \in \{0, \dots, 255\}$, $c_2 = \text{human_attention_seconds}$ (the makespan, §??), $c_3 = \text{token_cost}$, $c_4 = \text{latency_ms}$, and $c_5 = \text{context_switches} \in \{0, \dots, 255\}$ (distinct capabilities used). The order is lexicographic, cheapest-first, and the refused vector is `refused()` = (unadmitted, 255, ∞ , ∞ , ∞ , 255), the top element.

Theorem 8.1 (Lexicographic order is the infinite-weight limit). *Let the secondary coordinates be bounded, $|c_i| \leq M$ for $i \geq 1$ over the admitted plans (the risk and switch coordinates are ≤ 255 ; the makespan, latency, and token coordinates are bounded by the structural depth and duration constants of Theorem ??). Define the scalarization $C_\lambda(c) = \sum_{i=0}^5 \lambda^{5-i} c_i$ with $\lambda > 1$. Then there is a threshold Λ such that for all $\lambda > \Lambda$,*

$$c \preceq_{\text{lex}} c' \iff C_\lambda(c) \leq C_\lambda(c').$$

As $\lambda \rightarrow \infty$ the weight λ^5 on the admitted coordinate c_0 diverges relative to all others.

Proof. Let $c = (c_0, \dots, c_5)$ and $c' = (c'_0, \dots, c'_5)$ be two cost vectors. If $c = c'$, then $C_\lambda(c) = C_\lambda(c')$ holds immediately for any λ , and the equivalence is satisfied. If $c \prec_{\text{lex}} c'$, let $p \in \{0, \dots, 5\}$ be the first index at which the coordinates differ. By definition of the lexicographic order, we have $c_i = c'_i$ for all $i < p$, and $c'_p - c_p \geq \epsilon$, where $\epsilon > 0$ is the minimum coordinate resolution. We compute the difference in scalarized costs:

$$C_\lambda(c') - C_\lambda(c) = \sum_{i=0}^5 \lambda^{5-i} (c'_i - c_i) = \lambda^{5-p} (c'_p - c_p) + \sum_{i=p+1}^5 \lambda^{5-i} (c'_i - c_i).$$

Using the coordinate bound $|c'_i - c_i| \leq 2M$ for all $i \geq 1$ and the geometric sum formula, we lower bound the sum:

$$\sum_{i=p+1}^5 \lambda^{5-i}(c'_i - c_i) \geq -2M \sum_{i=p+1}^5 \lambda^{5-i} = -2M \frac{\lambda^{5-p} - 1}{\lambda - 1}.$$

Substituting this back into the difference yields:

$$C_\lambda(c') - C_\lambda(c) \geq \lambda^{5-p}\epsilon - 2M \frac{\lambda^{5-p} - 1}{\lambda - 1} > \lambda^{5-p} \left(\epsilon - \frac{2M}{\lambda - 1} \right).$$

For the difference to be strictly positive, it suffices that the term inside the parentheses is nonnegative:

$$\epsilon - \frac{2M}{\lambda - 1} > 0 \iff \lambda - 1 > \frac{2M}{\epsilon} \iff \lambda > 1 + \frac{2M}{\epsilon}.$$

Thus, choosing the threshold $\Lambda = 1 + 2M/\epsilon$, we guarantee that for all $\lambda > \Lambda$, the difference $C_\lambda(c') - C_\lambda(c)$ is strictly positive. Conversely, if $C_\lambda(c) \leq C_\lambda(c')$ for all $\lambda > \Lambda$, then we must have $c \preceq_{\text{lex}} c'$, since if $c' \prec_{\text{lex}} c$ we would have $C_\lambda(c') < C_\lambda(c)$ by the same argument, a contradiction. As $\lambda \rightarrow \infty$, the ratio of the weight of c_0 to any other coordinate weight λ^{5-i} for $i \geq 1$ is λ^i , which diverges to infinity, ensuring the absolute dominance of the admission coordinate. $\square$

Corollary 8.2 (Admitted dominates: no finite cost buys a refused route). *An admitted plan sorts strictly before any refused plan, regardless of its secondary costs: for any finite risk, makespan, token, latency, and switch values, an admitted `CostVector` is $\prec_{\text{lex}}$ the refused top element. Equivalently, the price of unlawfulness is infinite in the limit that defines the order.*

Proof. The lead coordinate c_0 is 0 for admitted and 1 for refused, and by Theorem ?? it dominates: any admitted vector differs from the refused element first at c_0 , where it is strictly smaller. Concretely, `CostVector`'s `Ord` places `admitted = true` ahead by reversing the boolean comparison, then breaks ties down the remaining coordinates; the `refused()` constructor sets c_0 to unadmitted and every secondary coordinate to its maximum (∞ for the reals, `u8::MAX/u64::MAX` for the integers), so it is the unique top. This is the test `cost_vector_orders_admitted_before_refused`. $\square$

Remark (The router's determinism is a receipt). `route_capability_plan` binds the problem text, the plan's steps, and the cost into a BLAKE3 `route_chain` [?]; identical tasks produce identical chains (`routing_is_deterministic_same_task_same_route`). The arbitration is thus not a heuristic tie-break but a receipted, replayable choice — the cost order is a function, and its value is committed.

8.1.1 Additive Separation of Maximum Reachable Revenue

Evaluating the global maximum of realizable revenue across accounts would normally result in exponential search complexity due to combinatorial action options.

Theorem 8.3 (Additive Separation of MRR). *Let A be a set of independent client accounts. Let $\text{realized}(a)$ be the revenue of account a under a target stage, and let $\text{lawful_targets}(a)$ be its evidence-gated target stages. The joint plan optimization decomposes linearly:*

$$\max_{\text{joint plan}} \sum_{a \in A} \text{realized}(a) = \sum_{a \in A} \max_{t \in \text{lawful_targets}(a)} \text{realized}(a, t). \quad (8.1)$$

Consequently, the search complexity is reduced from $\Omega(\prod_a |T_a|)$ to $O(|A| \cdot |T_a|)$ linear sum.

Proof. Let a joint plan be represented as a vector of target stages $\mathbf{t} = (t_a)_{a \in A}$, where $t_a \in \text{lawful_targets}(a)$ for each account $a \in A$. Since the client accounts are independent, the feasible choice set for each account a , denoted $\text{lawful_targets}(a)$, is independent of the choices made for all other accounts. Thus, the joint optimization domain $\mathcal{T}_{\text{joint}}$ is the Cartesian product of the individual feasible sets: $\mathcal{T}_{\text{joint}} = \prod_{a \in A} \text{lawful_targets}(a)$. The total realized revenue for a joint plan $\mathbf{t}$ is the sum of the individual revenues: $f(\mathbf{t}) = \sum_{a \in A} \text{realized}(a, t_a)$. Because the domain is a Cartesian product and the objective function $f(\mathbf{t})$ is additively separable, the optimization distributes over the sum:

$$\max_{\mathbf{t} \in \mathcal{T}_{\text{joint}}} \sum_{a \in A} \text{realized}(a, t_a) = \sum_{a \in A} \max_{t_a \in \text{lawful_targets}(a)} \text{realized}(a, t_a).$$

This proves equation (??). For the computational complexity, let $T_{\max} = \max_{a \in A} |\text{lawful_targets}(a)|$. The size of the joint search space is $\prod_{a \in A} |\text{lawful_targets}(a)|$, which is bounded by $O(T_{\max}^{|A|})$ and requires exponential search time if the choices were coupled. By separating the optimization, we solve $|A|$ independent maximum-search tasks, each checking at most $T_{\max}$ options. The total search cost is therefore the sum of the costs, which is bounded by $O(|A| \cdot T_{\max})$, establishing the linear complexity reduction. $\square$

8.2 The makespan functional and critical-path structure

The primary continuous cost coordinate c_2 is the makespan, and its structure is what `analyze_schedule` computes over the plan's POWL tape.

Definition 8.2 (Makespan functional; forward and backward pass). Given the precedence DAG of Construction ?? with per-op durations d_i and release times, the *forward pass* computes earliest starts and finishes

$$\text{es}_i = \max\left(r_i, \max_{j \prec i} \text{ef}_j\right), \quad \text{ef}_i = \text{es}_i + d_i,$$

the *makespan* is $T = \max_i \text{ef}_i$, and the *backward pass* computes latest starts $\text{ls}_i = \text{lf}_i - d_i$ with $\text{lf}_i = \min_{i \prec k} \text{ls}_k$ (or T at sinks). The *slack* of op i is $\text{ls}_i - \text{es}_i$, and the *critical path* is the set of zero-slack ops.

Proposition 8.4 (Makespan is the longest path; the critical path is zero-slack). *The makespan T equals the length of the longest ($\prec$ -)path in the precedence DAG weighted by durations, and an op has zero slack iff it lies on some longest path. Delaying a zero-slack op delays the makespan; delaying an op by less than its slack does not.*

Proof. The forward recursion $ef_i = es_i + d_i$ with $es_i = \max_{j \prec i} ef_j$ is the standard longest-path dynamic program over a DAG (topologically ordered here because op i 's predecessors have indices $< i$), so ef_i is the longest weighted path ending at i and $T = \max_i ef_i$ is the overall longest path. The backward pass gives the latest each op may start without increasing T ; slack $ls_i - es_i = 0$ iff i cannot move, i.e. it lies on a longest path. This is exactly `forward_pass`, `backward_pass`, and the `critical_path_mask` (bit i set iff slack $< 10^{-9}$) of `ScheduleAnalysis64`. $\square$

Proposition 8.5 (Capacity sensitivity is a finite difference of the found schedule). *Let $T(\phi = c)$ be the makespan of the schedule the planner finds when fluent ϕ 's initial capacity is c . The capacity delta reports $(T(c - 1), T(c), T(c + 1))$, and ϕ is binding iff $T(c - 1) \neq T(c)$ (or the problem is infeasible at $c - 1$). This is a finite-difference sensitivity of the specific schedule the greedy planner produced, not a dual shadow price of an optimal scheduler.*

Proof. `replan_with_perturbed_capacity` re-solves `find_temporal_plan` with ϕ 's initial value moved by ± 1 (clamped at 0) via `find_temporal_plan_with_fn_overrides`, and reads off the resulting makespan; `binding_resource_mask` sets bit i iff the -1 perturbation changed the makespan or made the problem infeasible. Because `find_temporal_plan` is a greedy forward-chaining planner (its own module comment disclaims optimality), the delta is the response of that found schedule, which is the honest claim the code makes — and no more. Infeasibility at $c - 1$ is itself evidence the resource binds, per Proposition ?? $\square$

Remark (The functional closes the loop). The makespan feeds back into arbitration: $c_2 = \text{human_attention_seconds} = \text{analysis.makespan}(\text{CostVector::from_analysis})$, so the primary continuous cost the lexicographic order minimizes is the longest-path length of Proposition ??. Attention conservation (Proposition ??) says the total attention spent is schedule-invariant; the makespan functional says only its *completion time* is what arbitration optimizes — exactly the keystone's "reordering redistributes when attention is spent, not how much."

Chapter 9

Conclusion

We gave the geometry beneath two of the enforced equation’s invariants. Mission manufacture is bounded search over a finite ground net (Theorem ??); the state it searches is a marking with a numeric fluent valuation, and the attention fluent, borrowed and returned, obeys a balance law whose nonnegativity is feasibility (Propositions ??–??). The reachable markings lie in a translated cone (Proposition ??), the safe-net specialization of which is the branchless firing rule the receipt path runs (Proposition ??). Conformance is membership in that cone, and a nonconformant trace is separated by a Farkas hyperplane — a bounded, sound certificate verifiable at $O(\kappa)$ (Theorem ??, Corollary ??). Token-replay fitness equals 1 exactly on genuine firing sequences, with the first disabled transition a localized witness (Theorem ??), and it is an honest population ratio (Proposition ??). POWL 2.0 gives the recursive partial orders and choice graphs (Definition ??); the plan-to-tape transitive reduction is the Hasse diagram of the precedence order (Proposition ??), and each node is judged in a working set bounded by its arity (Proposition ??). The Kourani–van der Aalst separable decomposition (Theorem ??) is the Rice quarantine for processes: separable is admitted, non-separable is refused with reason (Proposition ??). And cost arbitration is the infinite-weight limit in which the admitted coordinate dominates (Theorem ??, Corollary ??), over a makespan functional that is the longest path of the precedence DAG (Proposition ??).

The standard was never comprehension of the mission’s interior. It is a bounded search, a membership test with a small certificate, and a fitness value of one — each guaranteed by mechanism, each small enough to hold.

Conformance is distance to a cone; the certificate is a hyperplane.

Appendix A

Notation

Symbol	Meaning
M, M_0	marking; initial marking (vectors in $\mathbb{Z}_{\geq 0}^p$)
$\mathbf{m}_t^-, \mathbf{m}_t^+$	preset (consumed) / postset (produced) of transition t
N, x	incidence matrix; Parikh (firing-count) vector
$\mathcal{P}, \text{cone}(N)$	marking polytope; reachability cone
y	Farkas separating-hyperplane normal (nonconformance certificate)
φ	token-replay fitness $\in [0, 1]$ (Q16.16; 1.0 = 0x0001.0000)
$f_\phi(t)$	free level of numeric fluent ϕ at time t
$T, \text{es}_i, \text{ls}_i$	makespan; earliest / latest start of op i
$c, \preceq_{\text{lex}}$	cost vector; lexicographic cost order
$\mu, \mathcal{O}^*, \perp, \kappa$	manufacturing morphism; admitted space; refusal; verifier context

Appendix B

Code Correspondence

Every object above is anchored to the `bcinr-pddl` / `bcinr-powl` / `praxis` source.

Mathematical object	Code artifact	Location
Grounded temporal problem	<code>GroundTemporalProblem::build</code>	<code>bcinr-pddl/src/ground.rs</code>
Bounded grounding	<code>PDDL8_MAX_GROUND</code> , <code>BoundExceeded</code>	<code>bcinr-pddl/src/ground.rs</code>
Forward search (manufacture)	<code>GroundProblem::find_plan</code> (BFS)	<code>bcinr-pddl/src/ground.rs</code>
Numeric fluent update	<code>eval_numeric</code> , <code>apply_numeric_effect</code>	<code>bcinr-pddl/src/ground.rs</code>
Attention fluent	<code>(:functions (attention)) domain</code>	<code>bcinr-pddl/src/capability.rs</code>
Firing rule (safe net)	<code>PowlReplayVerifier::replay_frame</code>	<code>bcinr-powl-receipt/src/replay.rs</code>
Fitness / precision	<code>ConformanceMetrics</code> , <code>finalize</code>	<code>bcinr-powl-receipt/src/replay.rs</code>
Lifecycle token net	<code>TOK_*</code> , <code>replay_lifecycle</code>	<code>praxis-core/src/replay_adapter.rs</code>
POWL tape / reduction	<code>temporal_plan_to_powl_tape</code>	<code>bcinr-pddl/src/powl_bridge.rs</code>
Cost vector / order	<code>CostVector</code> , <code>impl Ord</code>	<code>bcinr-pddl/src/capability.rs</code>
Route receipt	<code>route_capability_plan</code>	<code>bcinr-pddl/src/capability.rs</code>
Makespan / critical path	<code>analyze_schedule</code> , <code>ScheduleAnalysis64</code>	<code>bcinr-pddl/src/schedule_analysis.rs</code>
Capacity sensitivity	<code>replan_with_perturbed_capacity</code>	<code>bcinr-pddl/src/schedule_analysis.rs</code>
CLI surface	<code>plan solve/analyze/route/execute</code>	<code>praxis/src/verbs/plan.rs</code>

Appendix C

The Datalog Index Layer: Nemo Saturation and Symbol Encoding

C.1 Datalog saturation as the reachability fixpoint

Forward planning (Chapter ??) searches over explicit markings. At the index layer, reachability can be decided faster by computing the *fixpoint* of a Datalog program that derives all entailed atoms from initial facts and action rules, before any search begins. The `pddl-index` crate embeds the Nemo Datalog engine to realize this.

Definition C.1 (Datalog saturation for grounded PDDL). Let $(\mathcal{D}, \mathcal{O}[])$ be a grounded PDDL problem with initial atom set M_0 and ground action set T . The *Nemo Datalog program* Π for this problem has one fact per atom in M_0 and one rule per ground action $t \in T$:

$$\text{add-effect}(t) \leftarrow \bigwedge_{\text{prec}} \text{prec}(t).$$

The *saturation* $\text{sat}(\Pi)$ is the least fixpoint $\text{lfp}(T_\Pi)$ of the immediate-consequence operator T_Π , which applies all rules until no new atom is derived. An atom p is *reachable* iff $p \in \text{sat}(\Pi)$.

Proposition C.1 (Saturation is sound and complete for forward reachability). $p \in \text{sat}(\Pi)$ iff there exists a sequence of ground actions whose add-effects include p , i.e. p is reachable from M_0 in the Petri-net sense of Chapter ??.

Proof. Soundness: every derived atom is the add-effect of a ground action whose preconditions were derived earlier (by induction on the derivation depth); the actions form a witness firing sequence. Completeness: a reachable atom has a witness sequence; unrolling it gives a derivation in T_Π . This is the standard Datalog fixpoint theorem [?, ?] applied to the grounded action rules. $\square$

Remark (Saturation prunes the BFS frontier). Once $\text{sat}(\Pi)$ is computed, the BFS of `find_plan` need not consider action t if any of its preconditions is absent from $\text{sat}(\Pi)$ — such an action cannot contribute to any plan reachable from M_0 . This is the `pddl-index` pruning step: the frozen fact store (Construction ??) is built from $\text{sat}(\Pi)$, and the XOR-filter membership gate discards unreachable preconditions before the search even begins.

C.2 Dictionary-encoded symbol identifiers

Definition C.2 (Symbol dictionary). A *symbol identifier* $\text{SymId} \in \mathbb{Z}_{>0}$ is a compact integer encoding of a ground atom or object name. The `pddl-index` crate maintains a bidirectional map between string symbols and SymId values:

$$\text{SymDict} : \{\text{strings}\} \leftrightarrow \mathbb{Z}_{>0},$$

encoded as a `FxHashMap<String, u32>` in the forward direction and a `Vec<String>` in the reverse (lookup by ID). The zero value is reserved as “undefined,” so all live IDs are strictly positive.

Proposition C.2 (Dictionary encoding reduces grounding cost). *With symbol dictionary encoding, a ground atom $p(o_1, \dots, o_k)$ is represented as a tuple of $k + 1$ SymId values (predicate ID plus argument IDs), each a 32-bit integer. Equality testing reduces to integer comparison; hash-set membership (e.g. for the initial atom set M_0) is a single hash over $4(k + 1)$ bytes. The cost of the grounding loop (Theorem ??) with encoding is $O(|T| \cdot k \cdot N)$ integer operations, replacing $O(|T| \cdot k \cdot N)$ string comparisons with their $O(1)$ -time integer counterparts.*

Proof. Each parameter binding substitutes an object’s SymId (an integer lookup in a `Vec`, $O(1)$) for a parameter name. Forming the atom tuple is $O(k)$ integer writes. The hash of a tuple of 32-bit integers is computed by `FxHash` in $O(k)$ time independent of string lengths. Summing over $|T|$ schemas and N^k bindings gives the stated cost. $\square$

Construction C.1 (Relational join for type-safe grounding). Grounding with type constraints is a relational join: let $\mathcal{T}$ be the type hierarchy (Definition ??) encoded as a parent array `parent: Vec<SymId>`. For a schema parameter v with type τ , the admissible objects are $\{o : \text{TypeIndex}::\text{satisfies}(\text{SymId}(o), \text{SymId}(\tau))\}$, where `satisfies` walks the parent chain upward until either $\text{SymId}(\tau)$ is found (admissible) or the root is reached (not admissible). This walk has depth $\leq$ the height of $\mathcal{T}$, a design constant independent of $|\mathcal{O}|$. The join loop iterates only over type-compatible objects for each parameter, reducing the effective branching factor from N^k to $\prod_i |\{o : \text{type}(o) \leq \tau_i\}|$, which for narrow types is much smaller.

Remark (Symbol encoding and the XOR filter compose). Construction ?? produces SymId tuples; the XOR filter (Construction ??) applies FNV-1a to those tuples. Since FNV-1a is a deterministic function of its input bytes and SymId values are fixed-width 32-bit integers, the filter is canonical across runs: rebuilding the fact store from the same saturation produces the same filter bytes, enabling content-addressable caching of the grounded index.

Bibliography

- [1] R. A. Kowalski, *Predicate logic as programming language*, Information Processing **74** (Proc. IFIP Congress), North-Holland, 1974, 569–574.
- [2] J. W. Lloyd, *Foundations of Logic Programming*, 2nd ed., Springer-Verlag, 1987.
- [3] S. Chatman, *The Chatman Equation and the Industrial Revolution of Knowledge: $A = \mu(O)$, Knowledge Hooks, and Production-Verified Enterprise Execution*, v1.2.0, November 2025.
- [4] M. Fox and D. Long, *PDDL2.1: An extension to PDDL for expressing temporal planning domains*, Journal of Artificial Intelligence Research **20** (2003), 61–124.
- [5] T. Murata, *Petri nets: Properties, analysis and applications*, Proceedings of the IEEE **77**(4) (1989), 541–580.
- [6] A. Schrijver, *Theory of Linear and Integer Programming*, Wiley, 1986.
- [7] A. Rozinat and W. M. P. van der Aalst, *Conformance checking of processes based on monitoring real behavior*, Information Systems **33**(1) (2008), 64–95.
- [8] H. Kourani and S. J. van Zelst, *POWL: Partially Ordered Workflow Language*, in Business Process Management (BPM 2023), Lecture Notes in Computer Science **14159**, Springer, 2023, 92–108.
- [9] W. M. P. van der Aalst, *Process Mining: Data Science in Action*, 2nd ed., Springer, 2016.
- [10] W. M. P. van der Aalst, A. H. M. ter Hofstede, B. Kiepuszewski, and A. P. Barros, *Workflow patterns*, Distributed and Parallel Databases **14**(1) (2003), 5–51.
- [11] J. O’Connor, J.-P. Aumasson, S. Neves, and Z. Wilcox-O’Hearn, *BLAKE3: One function, fast everywhere*, 2020.